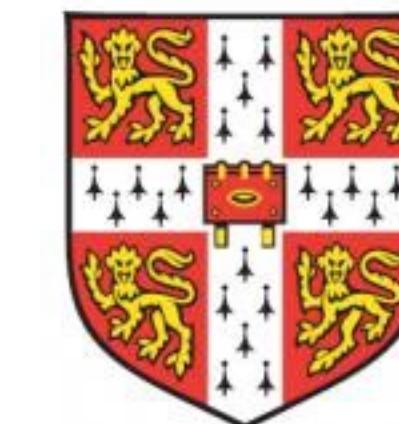
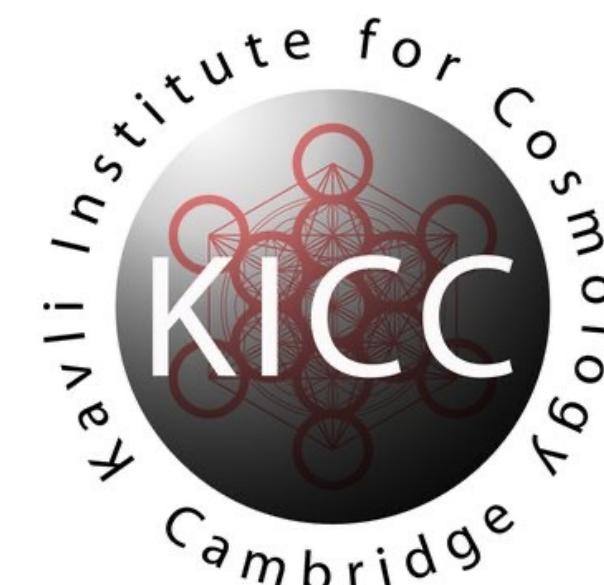


# **Diffusion meets Nested Sampling**

**Hills Coffee Talks**

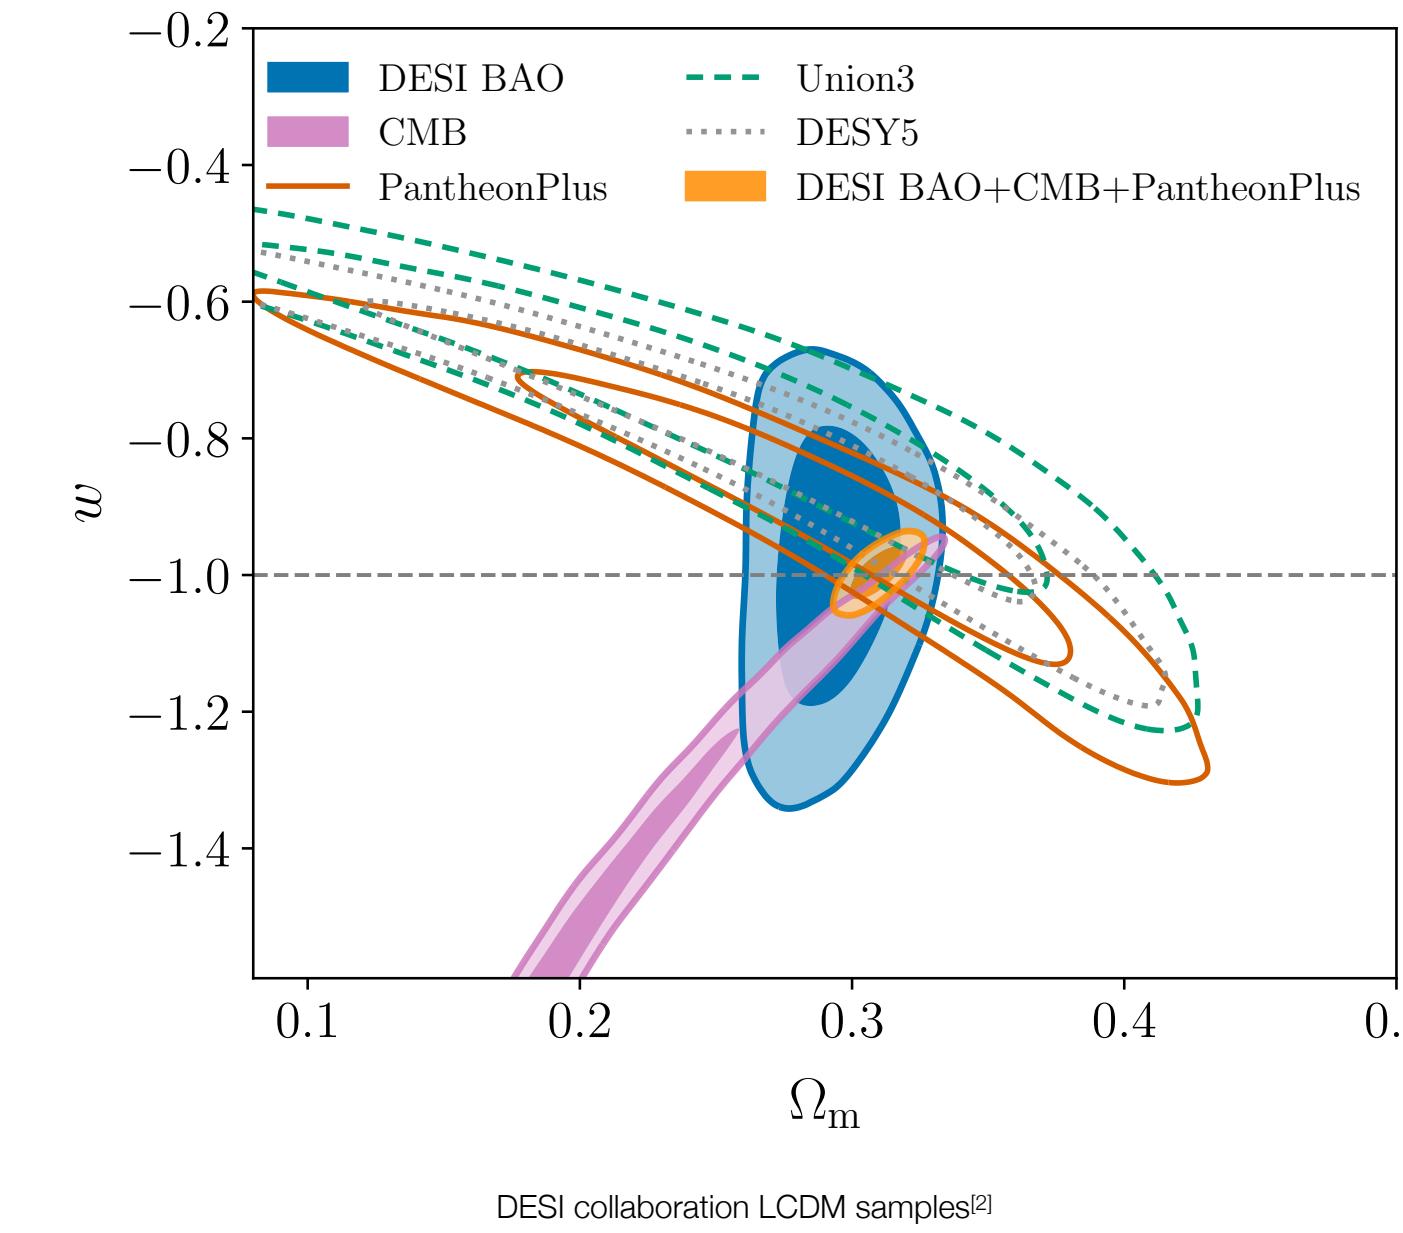
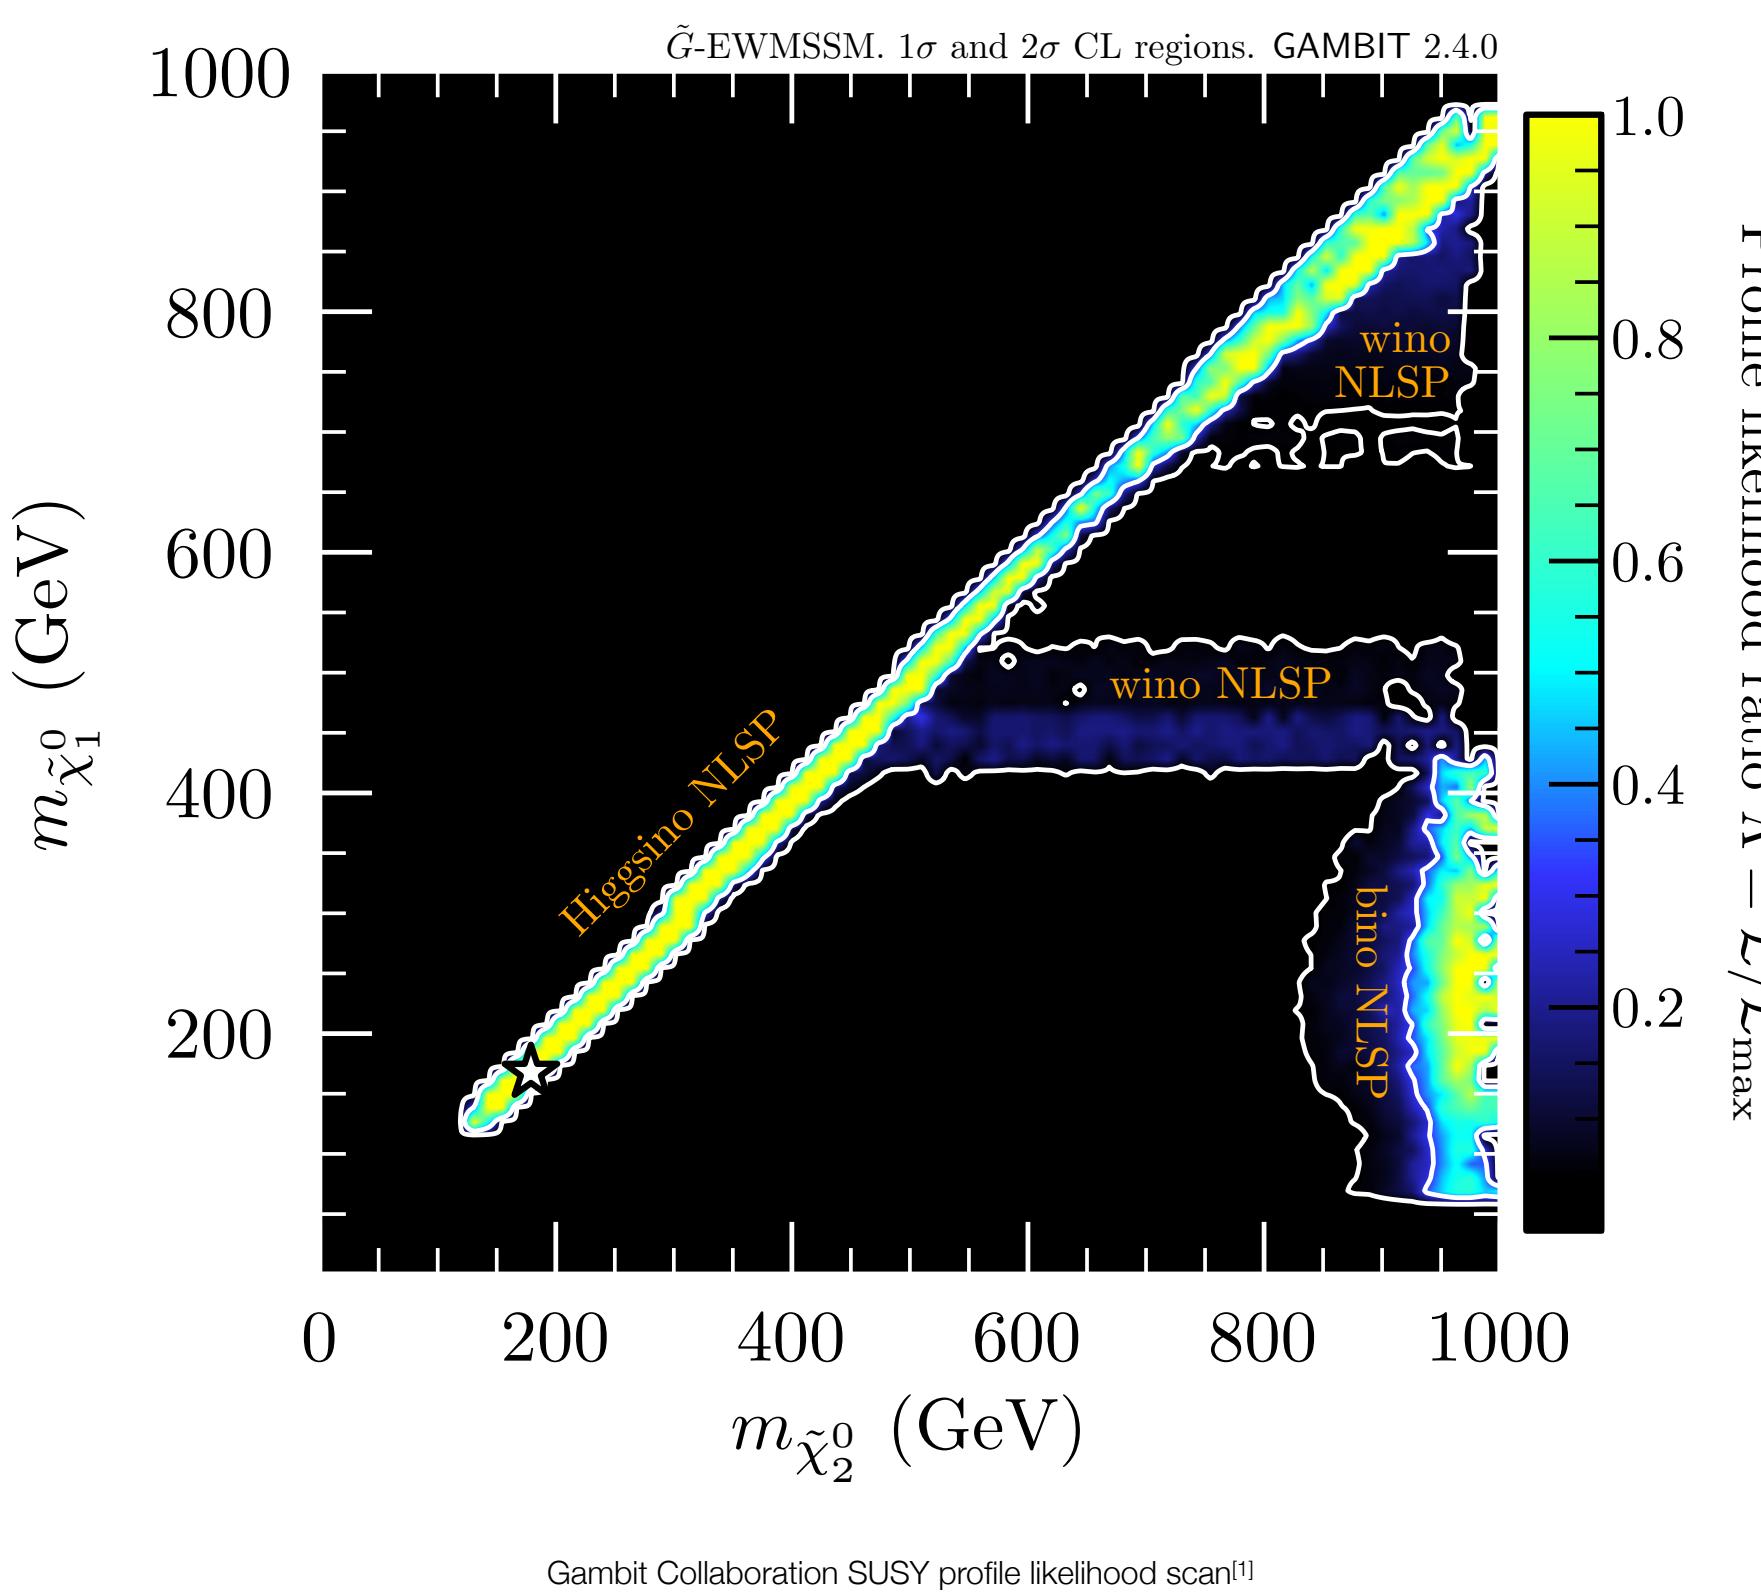


**UNIVERSITY OF  
CAMBRIDGE**

**David Yallup - 21/05**

# INFERENCE

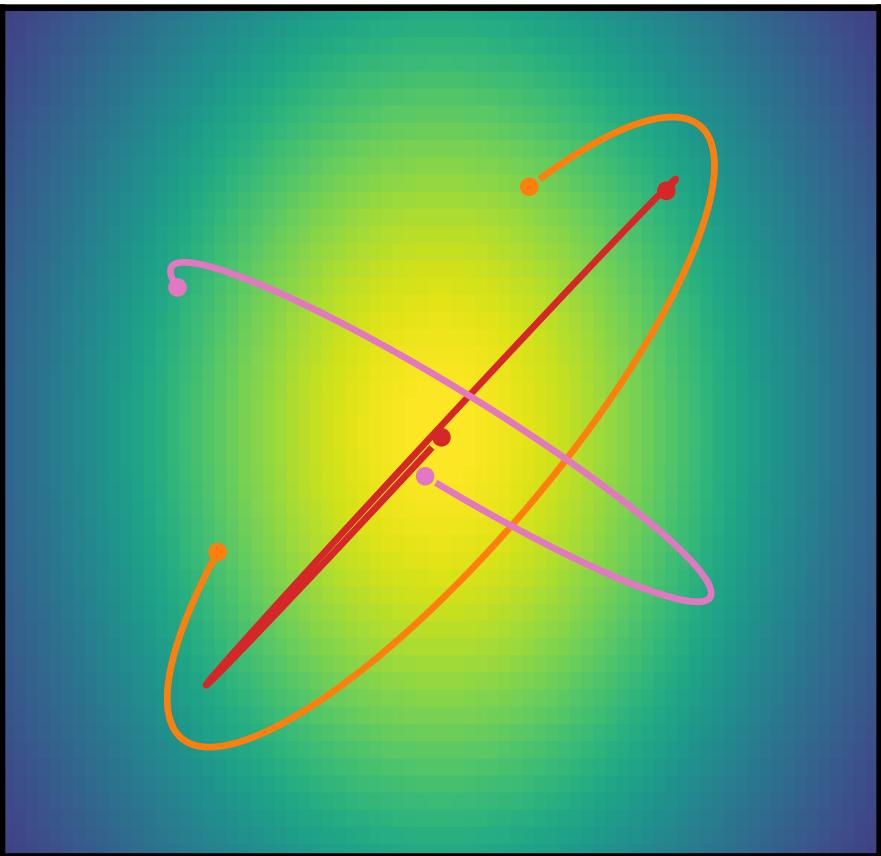
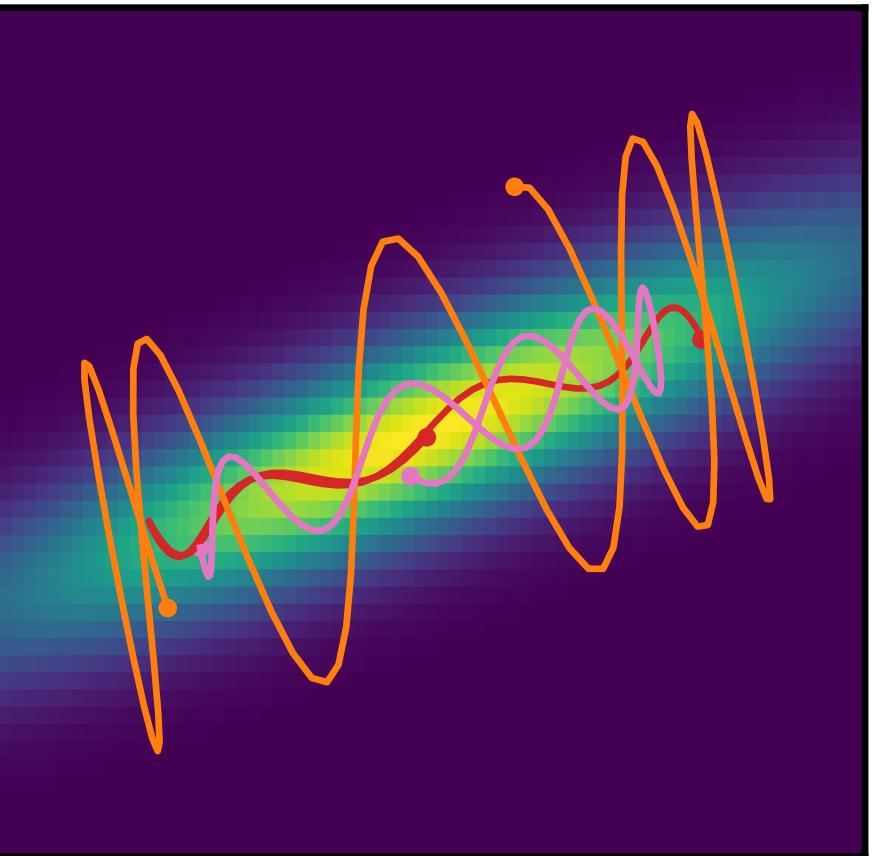
Fundamental physics is full of hard inference problems. Our optimization or sampling algorithms have to be able to navigate complex geometry



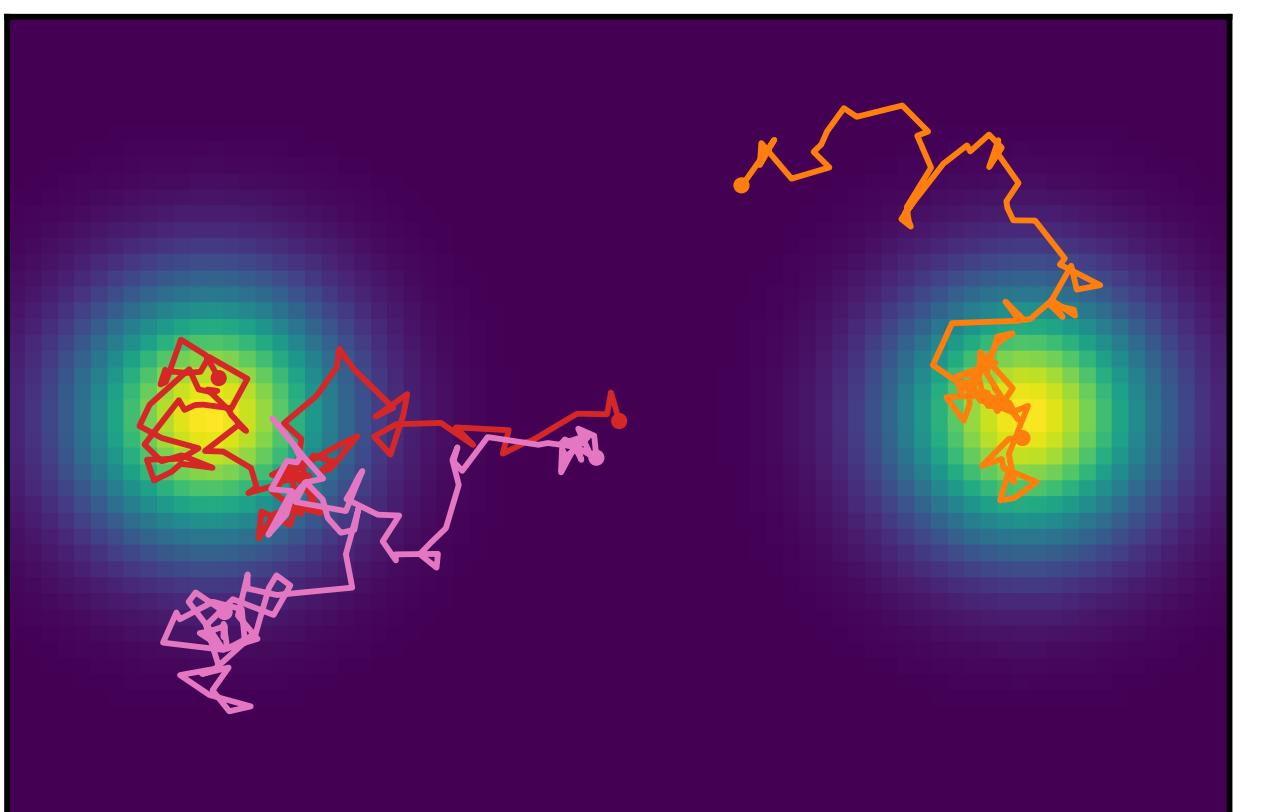
# GEOMETRY

Bad geometry<sup>[3]</sup> in inference problems comes in many guises, and intuition gets progressively less clear in high dimension. Machine learnt neural mappings offer us a new tool to approach this.

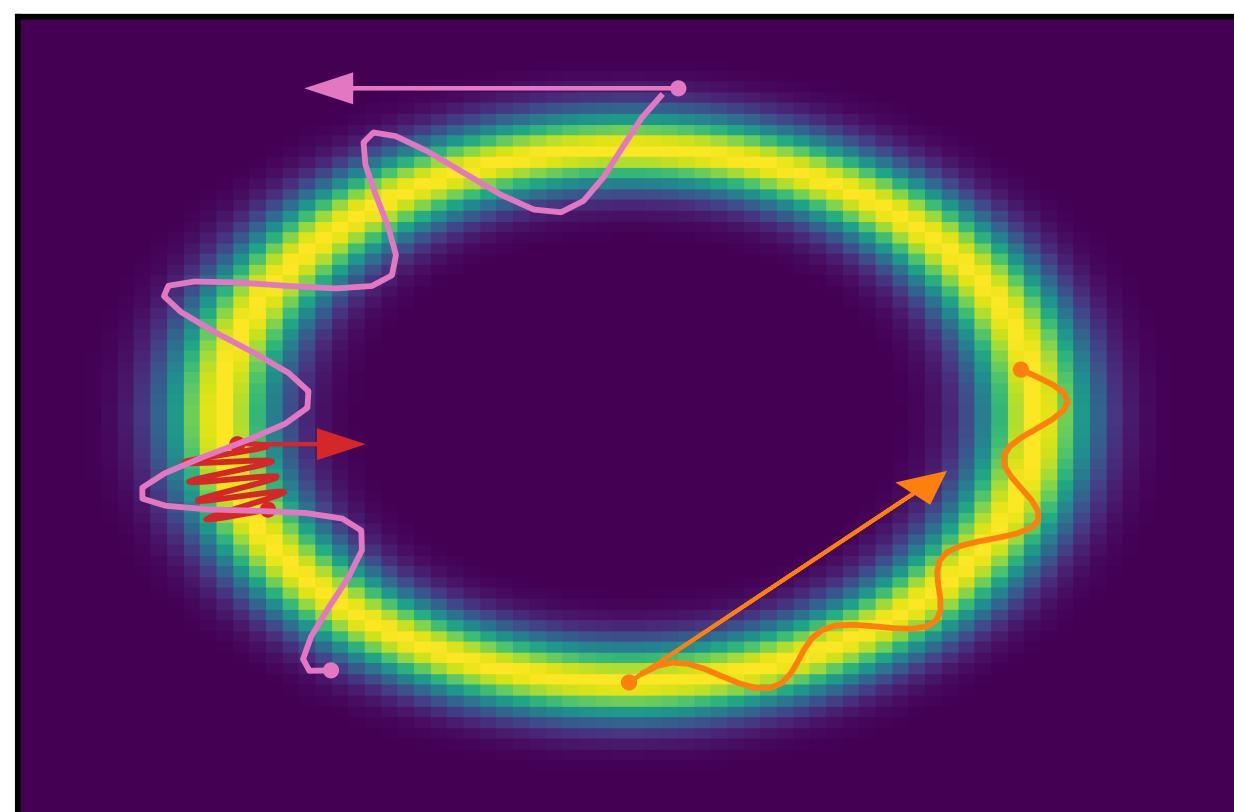
In the context of bridging distributions see pocoMC<sup>[4]</sup>, nessai<sup>[5]</sup>



Whitening transforms to regularize the metric.



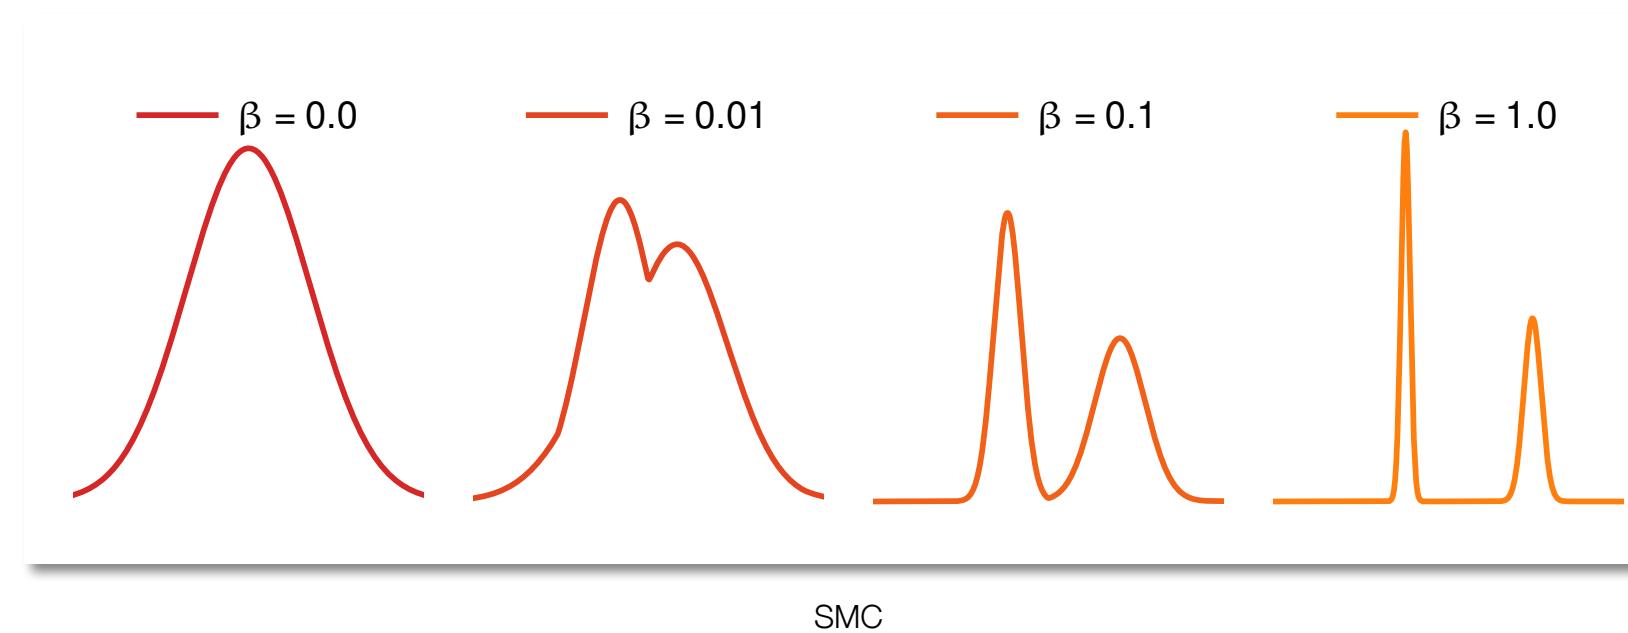
Clustering/ensembling to deal with multimodalities.



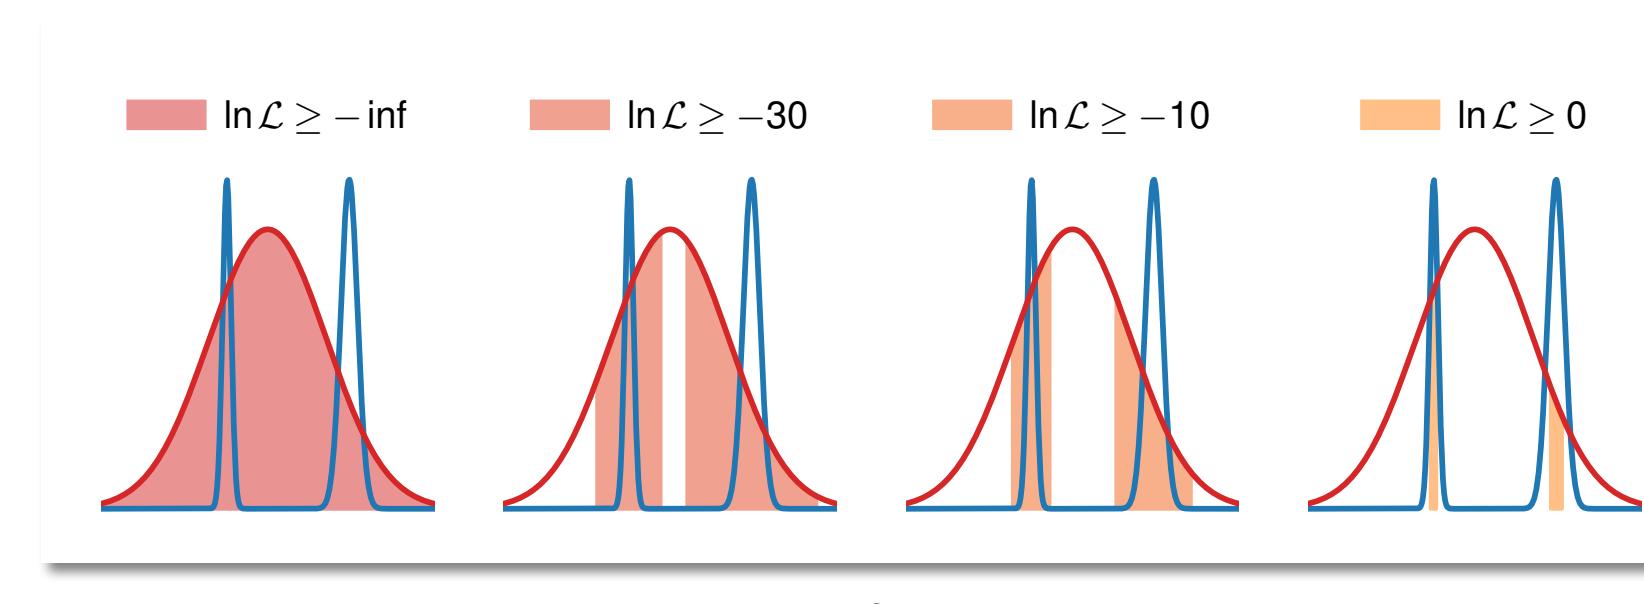
Gradients efficiently explore sweeping degeneracies.

# BRIDGING DISTRIBUTIONS

Population Monte Carlo methods — particle filters — form bridges from known (prior) to complex unknown (posterior) distributions. Sequential Monte Carlo (SMC) and Nested Sampling (NS) are two variants evolving populations of points<sup>[6]</sup>. Both give us access to the normalizing constant  $Z$ .



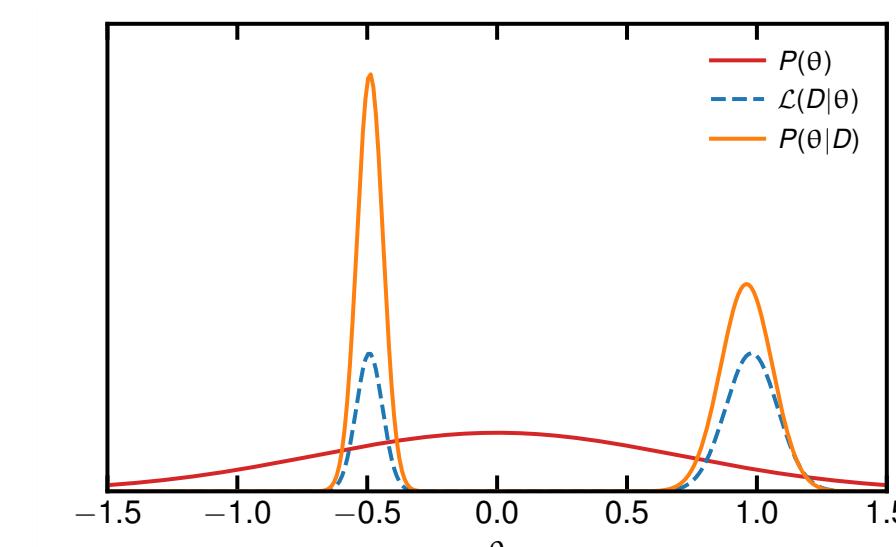
SMC



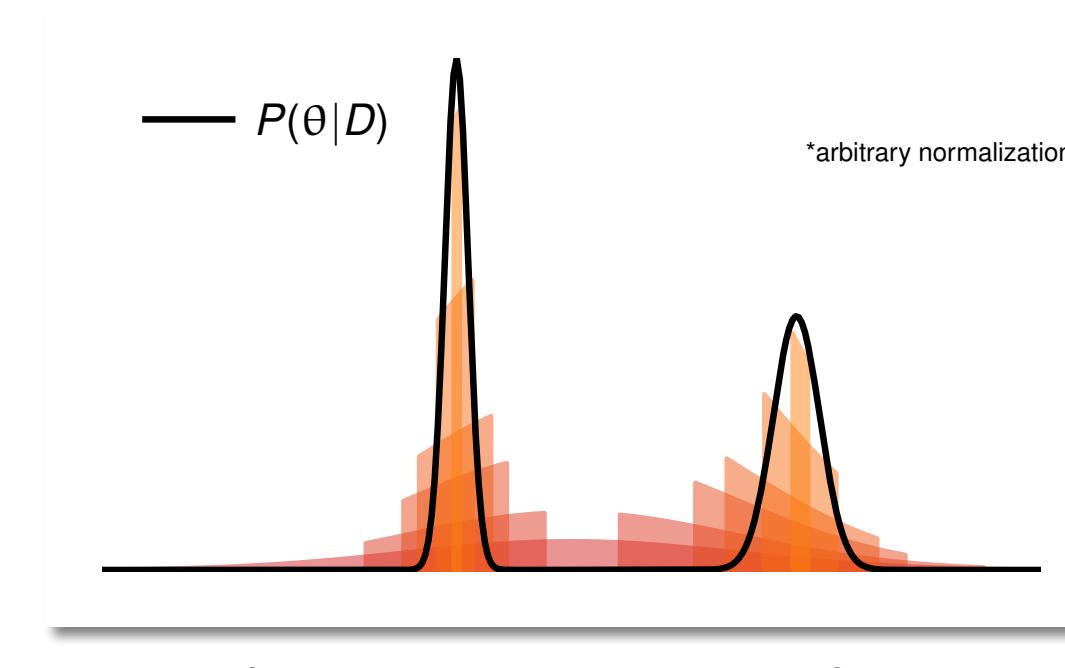
NS

$$P(\theta | D) = \frac{\mathcal{L}(D | \theta)P(\theta)}{Z}$$

Bayes Rule



Toy 1D inference problem.

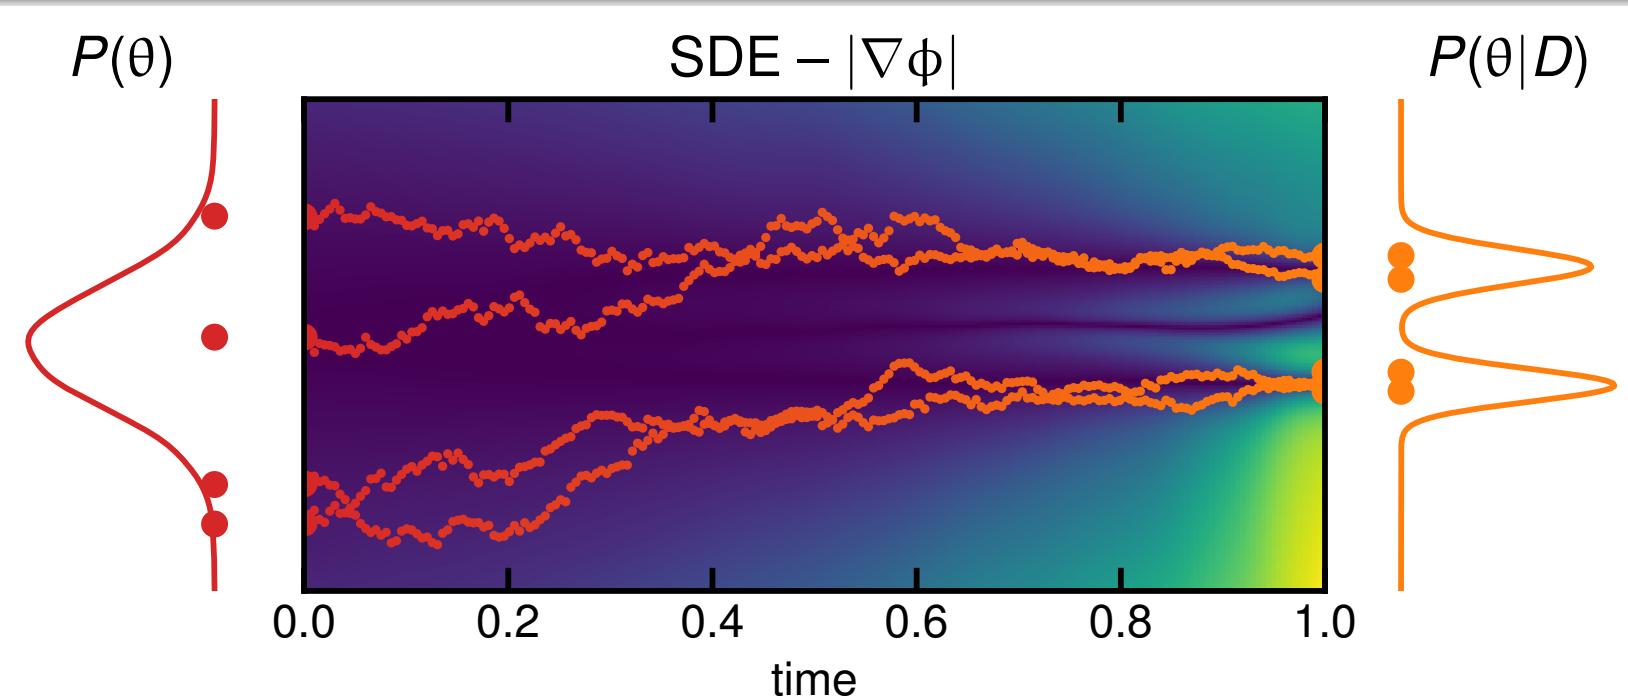
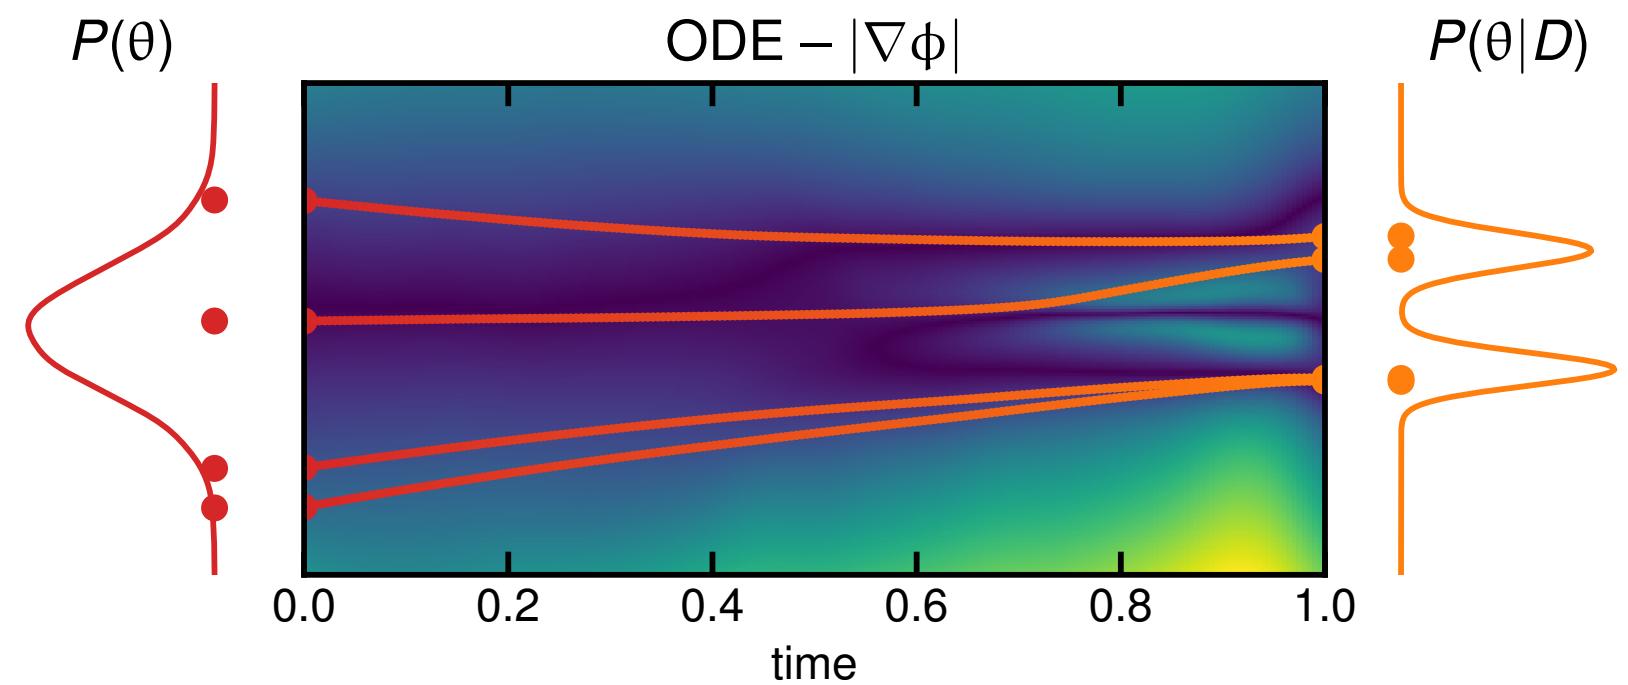


Schematic reconstructing posterior from NS shells.

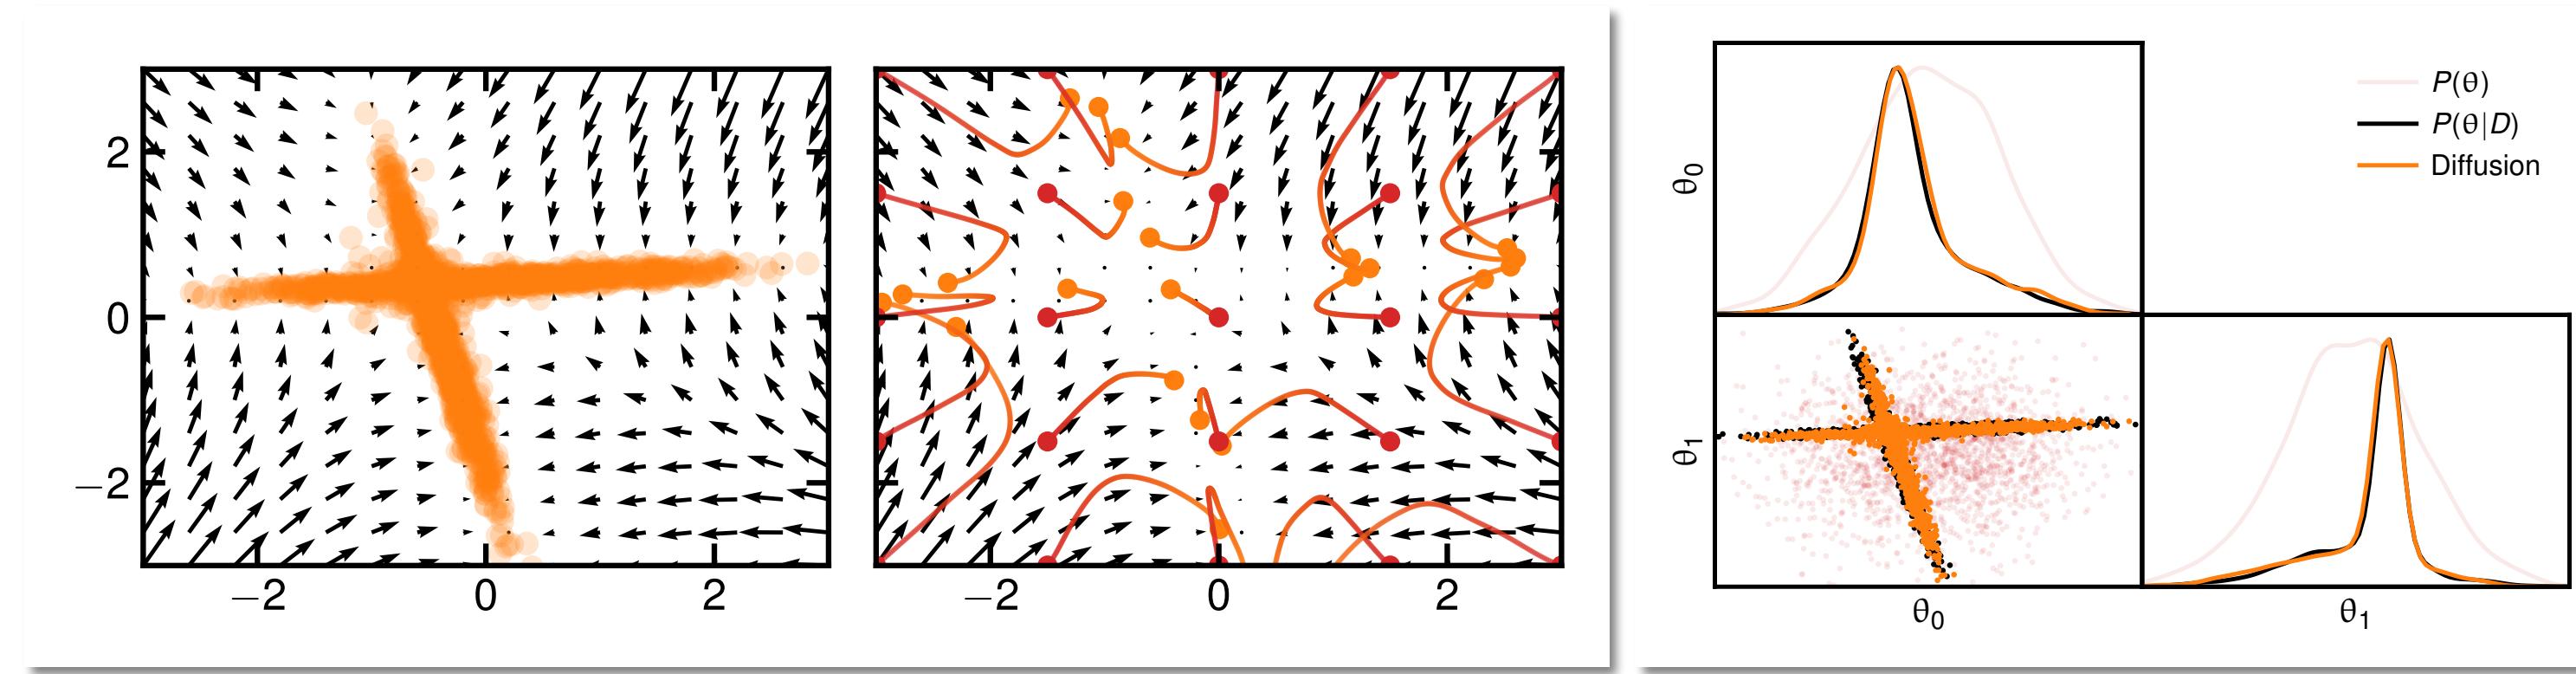
# DIFFUSION

Diffusion models learn the gradient of the implicit density of a point cloud. Solving evolution through this field with Stochastic Differential Equation (SDE) or Ordinary Differential Equation (ODE) solvers yields Diffusion<sup>[7]</sup> or Continuous flows<sup>[8]</sup>.

Neural learnt maps can transport any known distribution to an implicit target, no strict requirement on latent/prior!



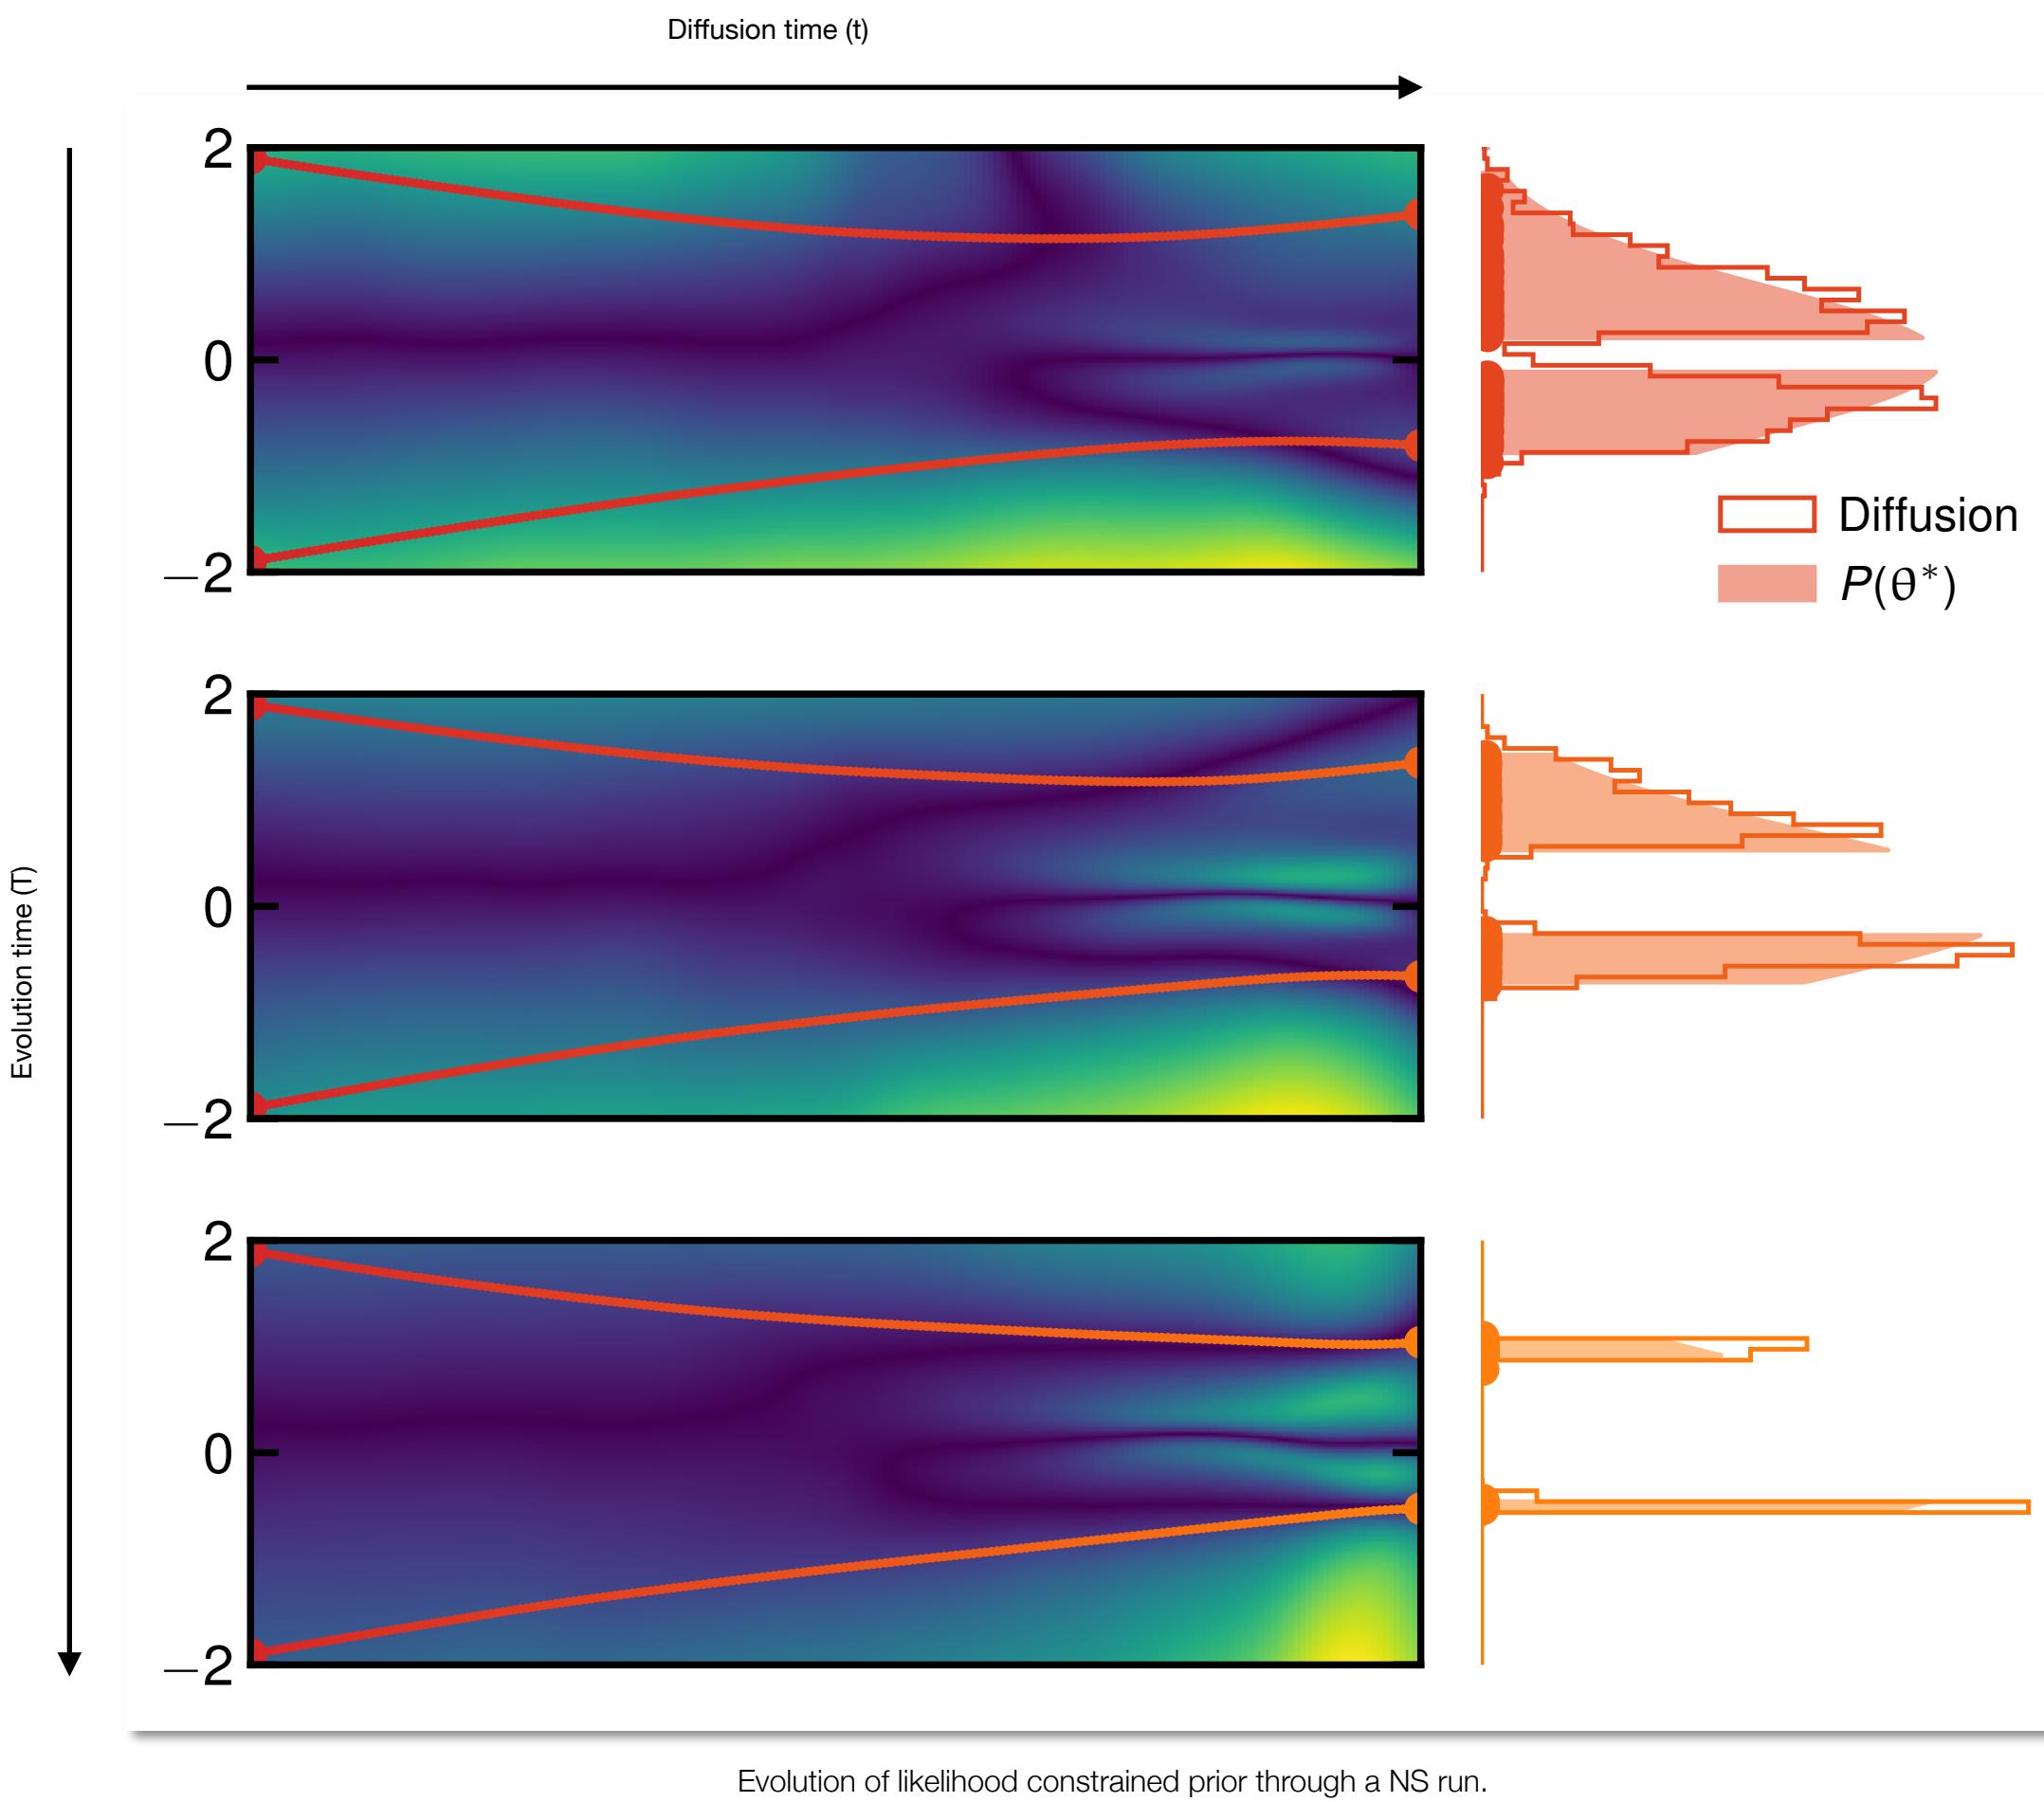
Learnt gradient vector field mapping prior to posterior.



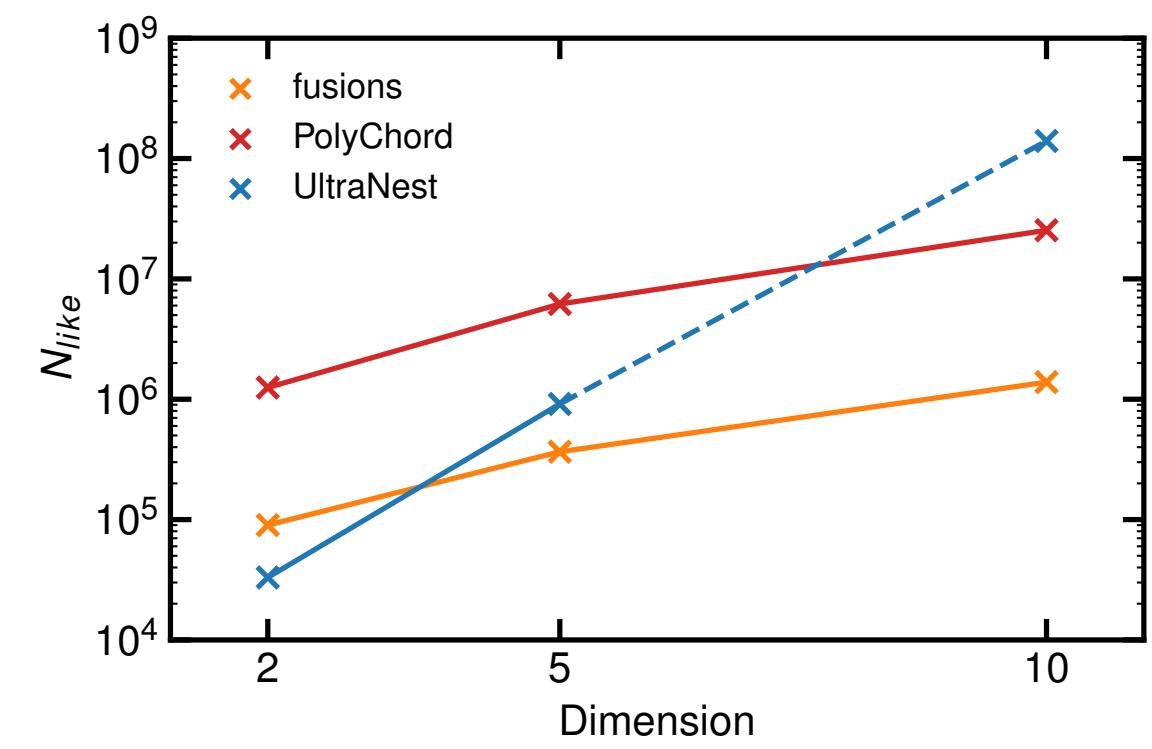
2D representation of learnt vector fields.

# RESULTS\*

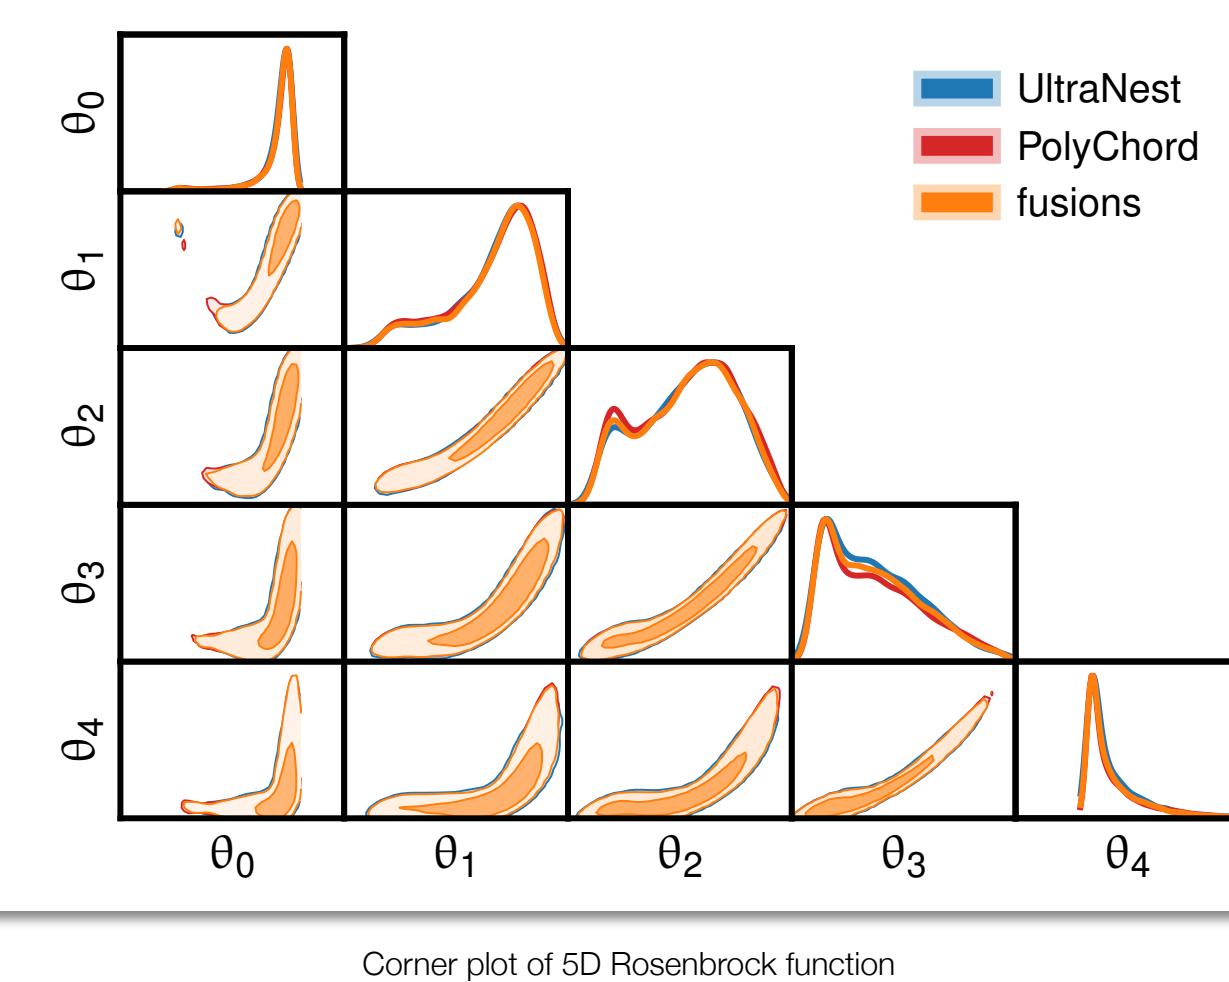
Diffusion models introduce time axis to the problem, bridging algorithms have another time axis we can efficiently evolve by fine tuning the score estimate.



# RESULTS\*

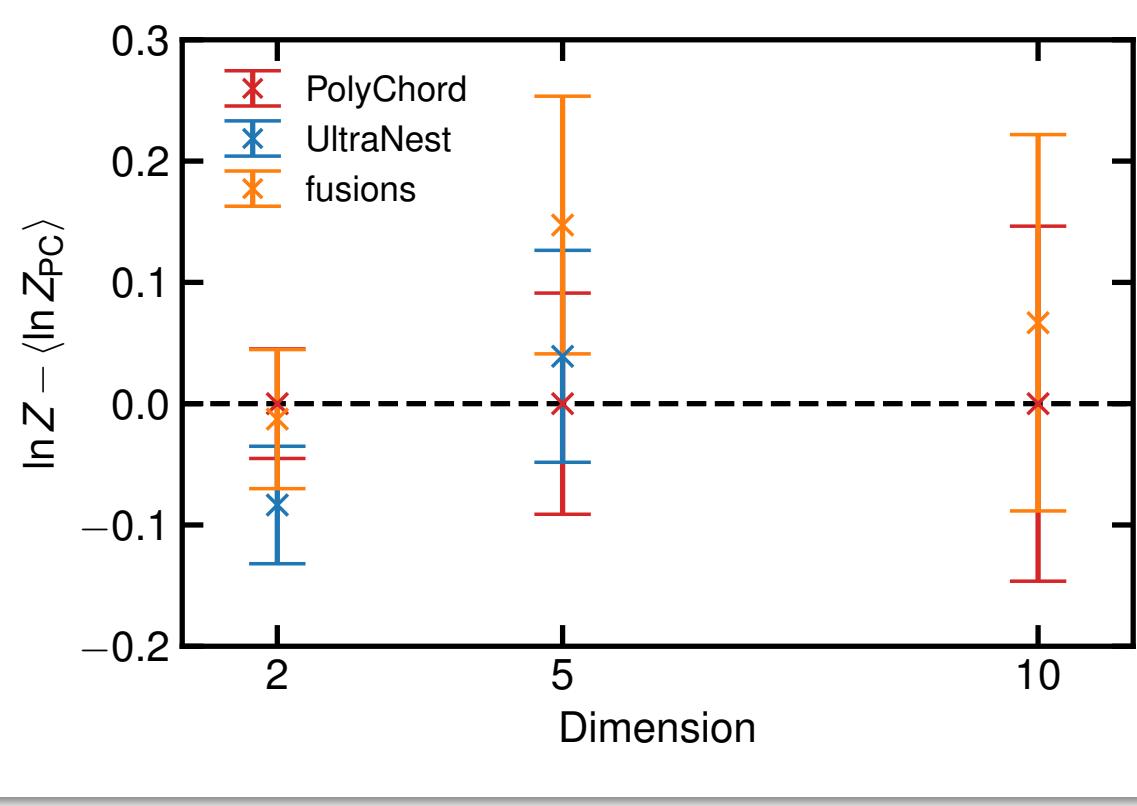
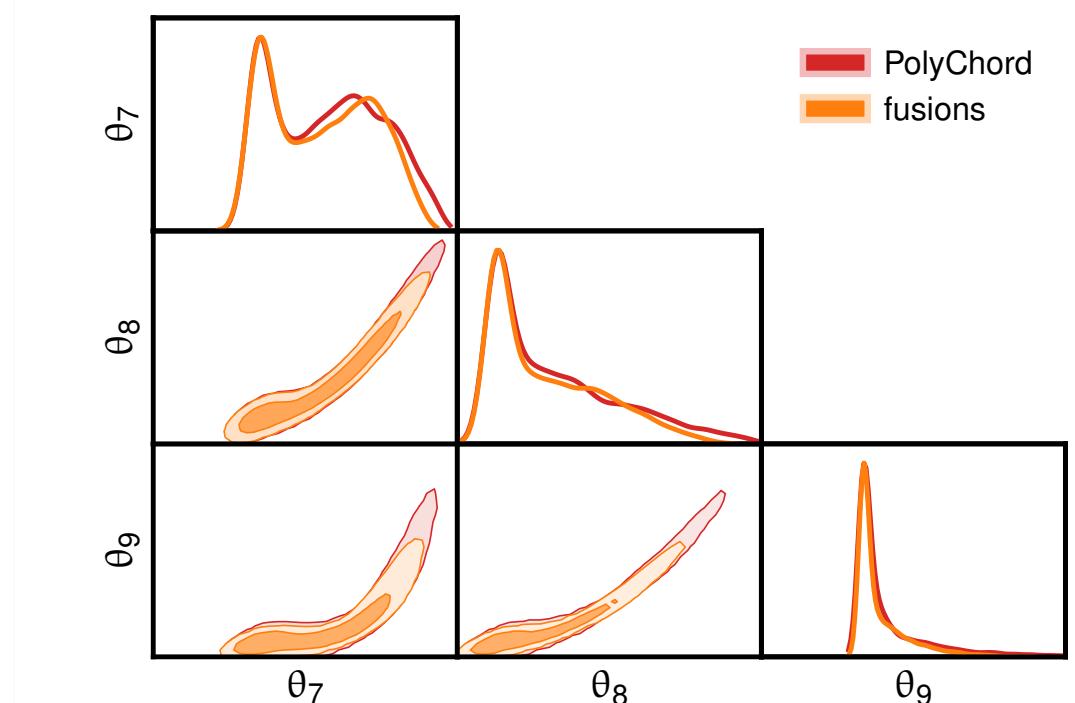


Number of function evaluations required for Rosenbrock function in various dimensions



$N_{\text{live}} = 2000$  for all, otherwise following defaults.

UltraNest in 10D  $N_{\text{like}}$  projected after early termination due to exceeding walltime.



Calculated log integral for Rosenbrock function in various dimensions

Comparison to standard (non-neural) tools<sup>[9,10]</sup> shows promising scaling, comparable to step samplers despite using rejection sampling, whilst maintaining accurate predictions on benchmark challenging problems.

Algorithm demonstrated uses zero classical methods, treating the geometry of the problem solely with neural networks and score based models.

\* Work in progress, comparison to other neural methods<sup>[4,5,11,12]</sup>, plenty left on the table to tune in the algorithm.

# DIFFUSION MEETS NESTED SAMPLING

## NEUTRALISING BAD GEOMETRY IN BRIDGING INFERENCE PROBLEMS

DAVID YALLUP



yallup/fusions



dy297@cam.ac.uk



yallup@github.io

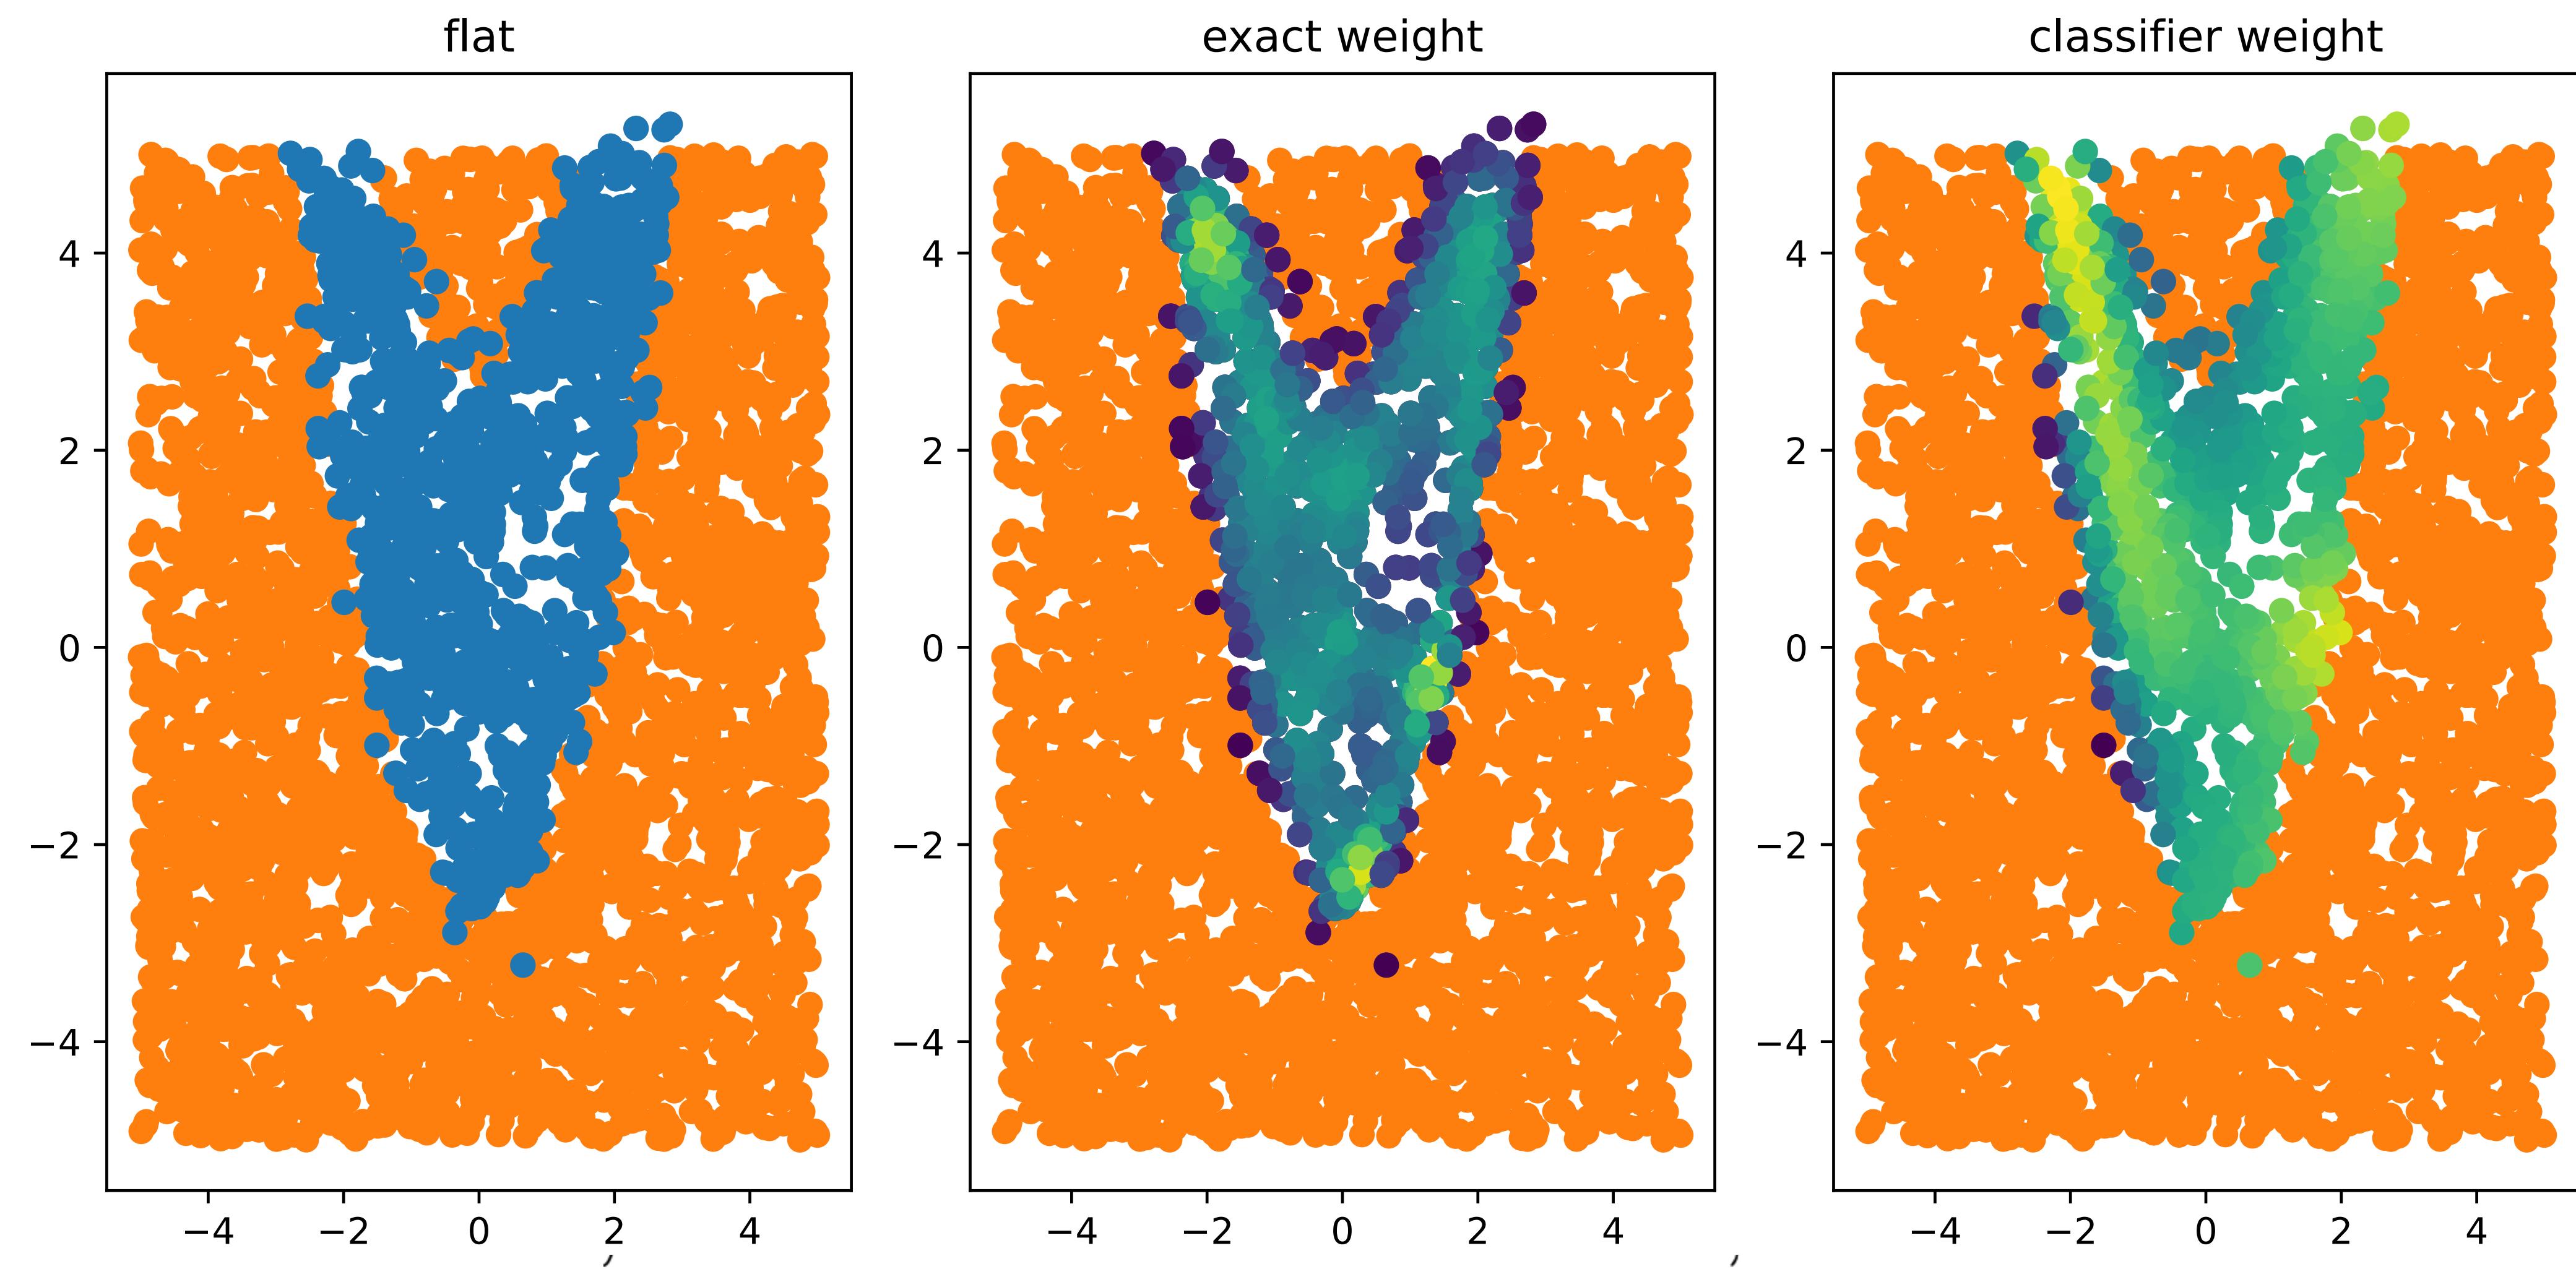
References:

- 1.[2303.09082] The Gambit collaboration
- 2.[2404.03002] DESI Collaboration
- 3.[1903.03704] Hoffman et al.
- 4.[2207.05660] Karamanis et al.
- 5.[2102.11056] Williams et al.

- 6.[2205.15570] Ashton et al.
- 7.[2011.13456] Song et al.
- 8.[2302.00482] Tong et al.
- 9.[2101.09604] Buchner
- 10.[1506.00171] Handley et al.
- 11.[2306.16923] Lange
- 12.[1903.10860] Moss

Technical references:

- [github.com/patrick-kidger/diffrax](https://github.com/patrick-kidger/diffrax)
- [github.com/handley-lab/anesthetic](https://github.com/handley-lab/anesthetic)
- [github.com/yallup/fusions](https://github.com/yallup/fusions)
- [github.com/google/flax](https://github.com/google/flax)
- [github.com/google/jax](https://github.com/google/jax)

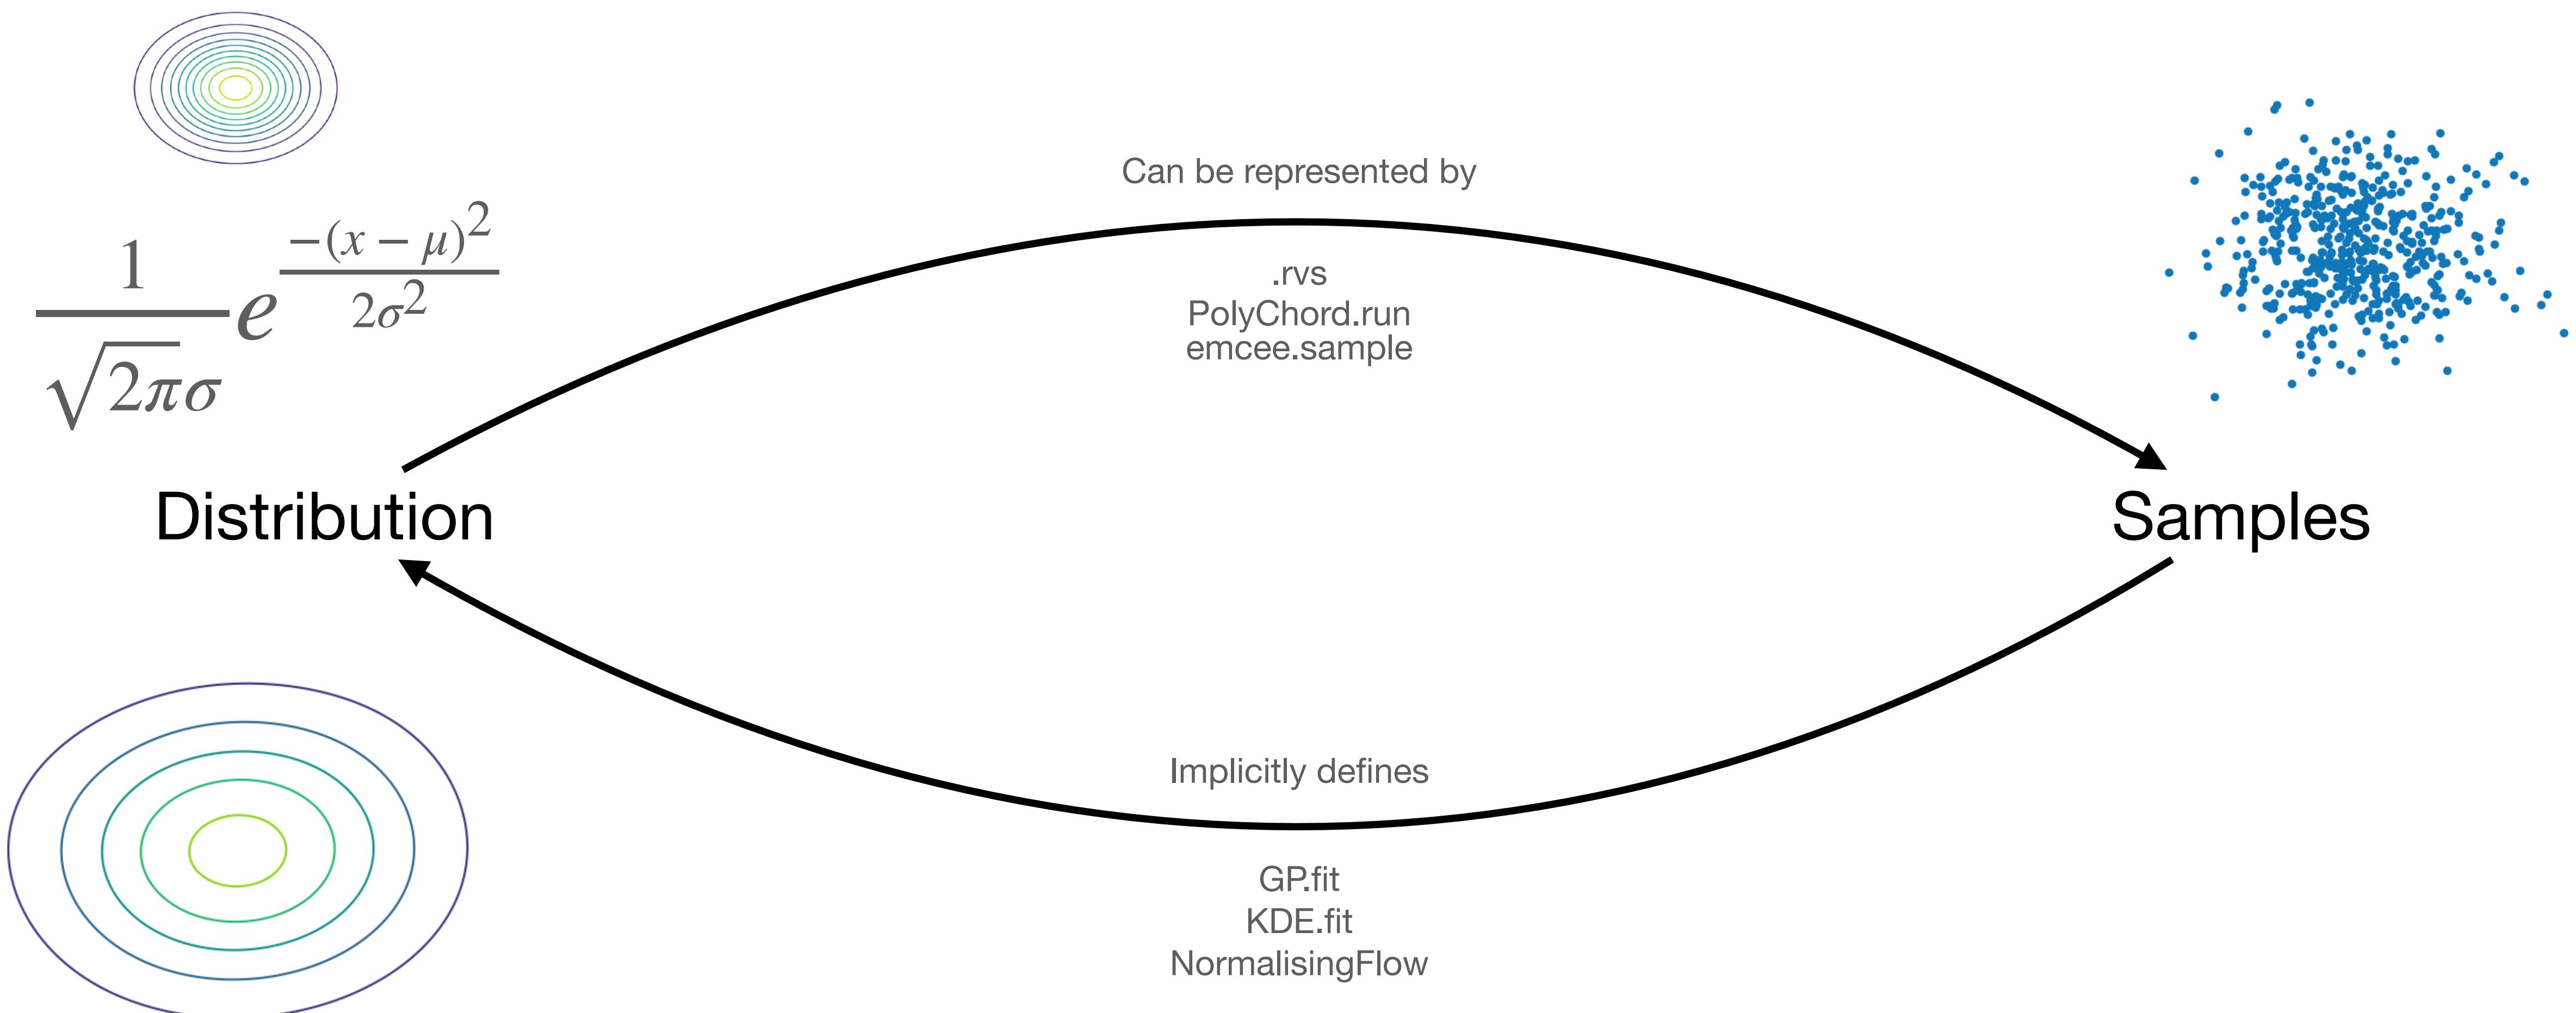


$$\nabla \log p_t(x_t|y) = \underbrace{\nabla \log p_t(x_t)}_{\text{pre-trained score}} + \underbrace{\nabla \log c_t(y|x_t)}_{\text{guidance}},$$

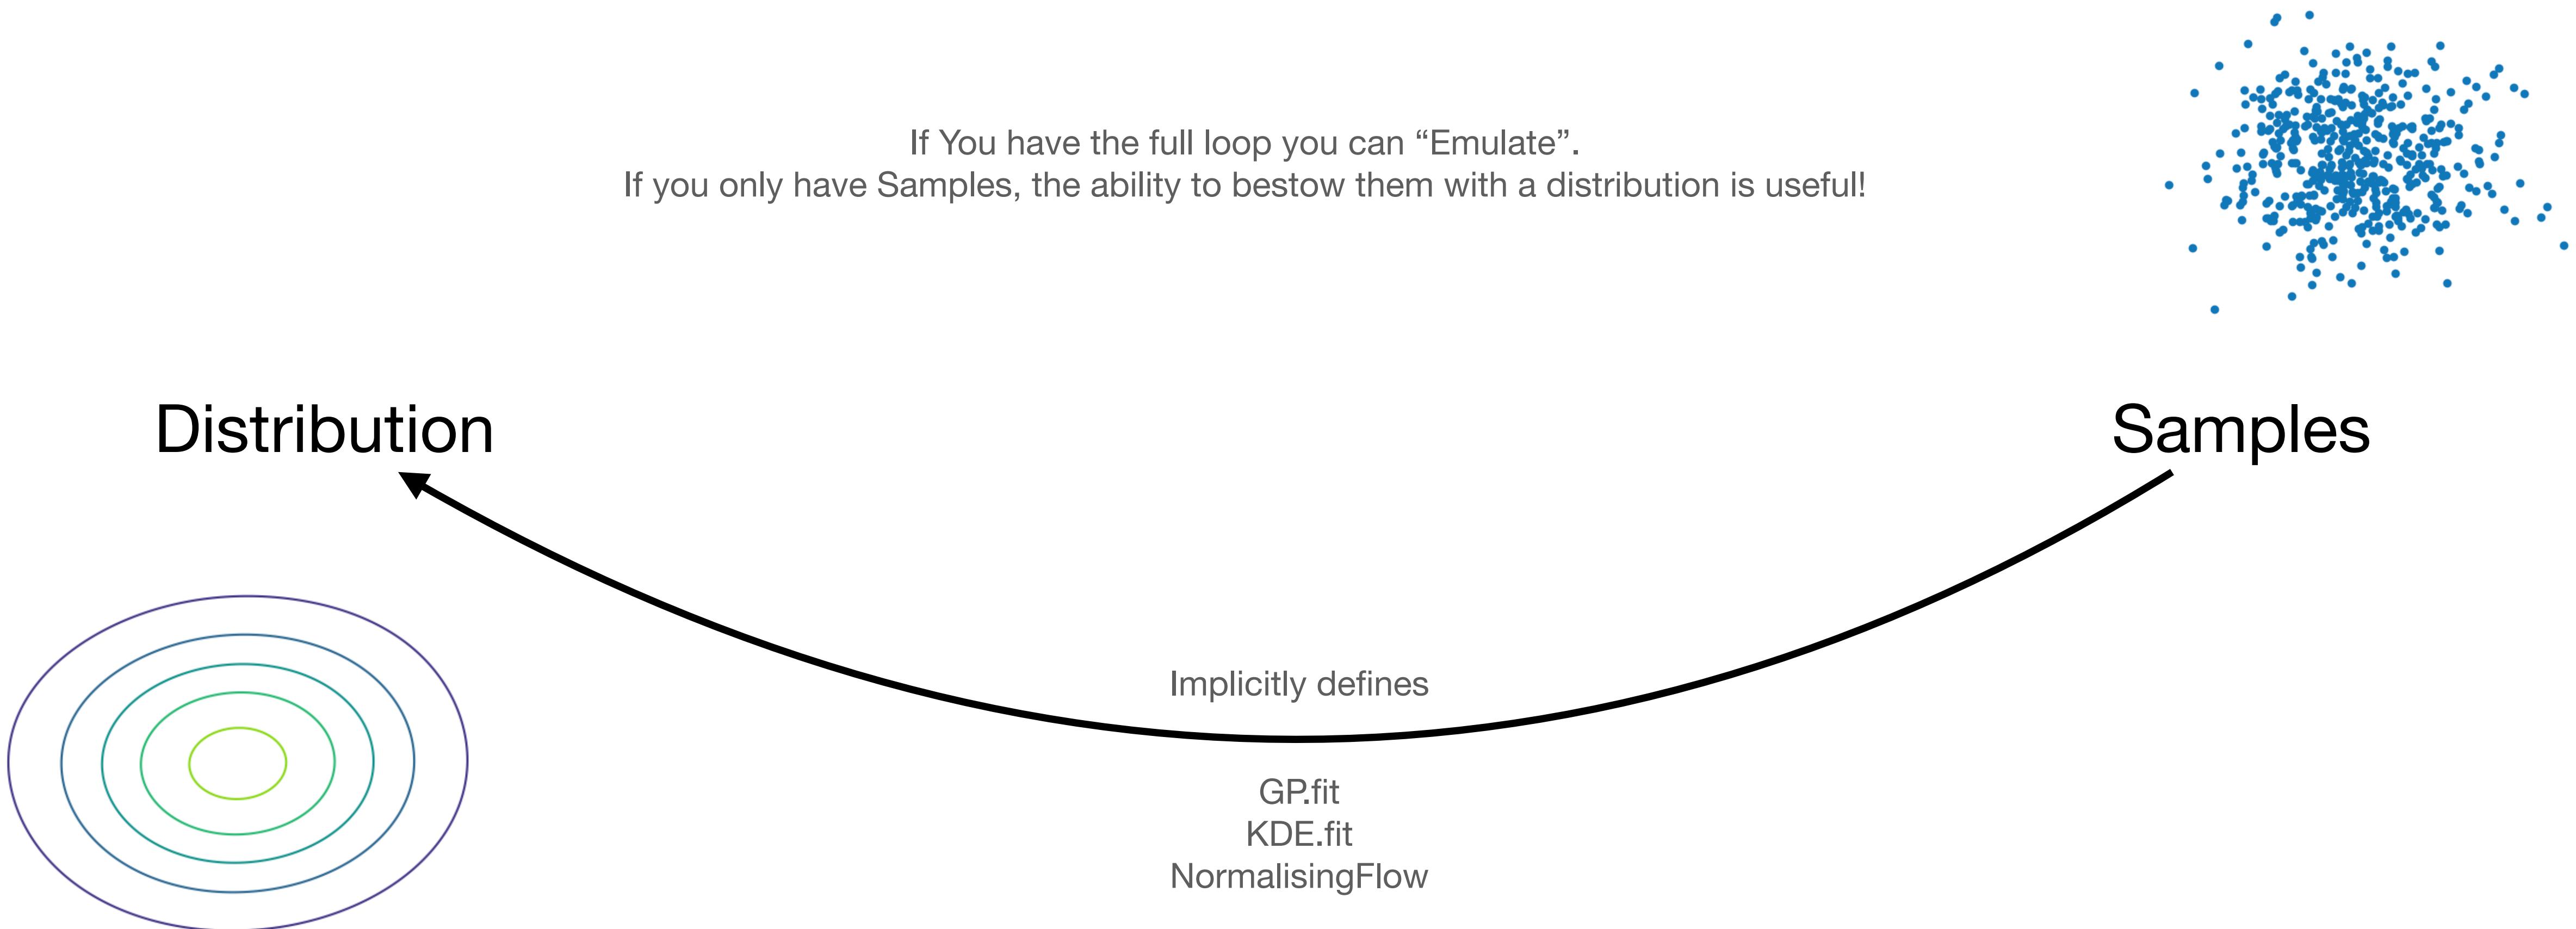


# Backup

# Modelling distributions

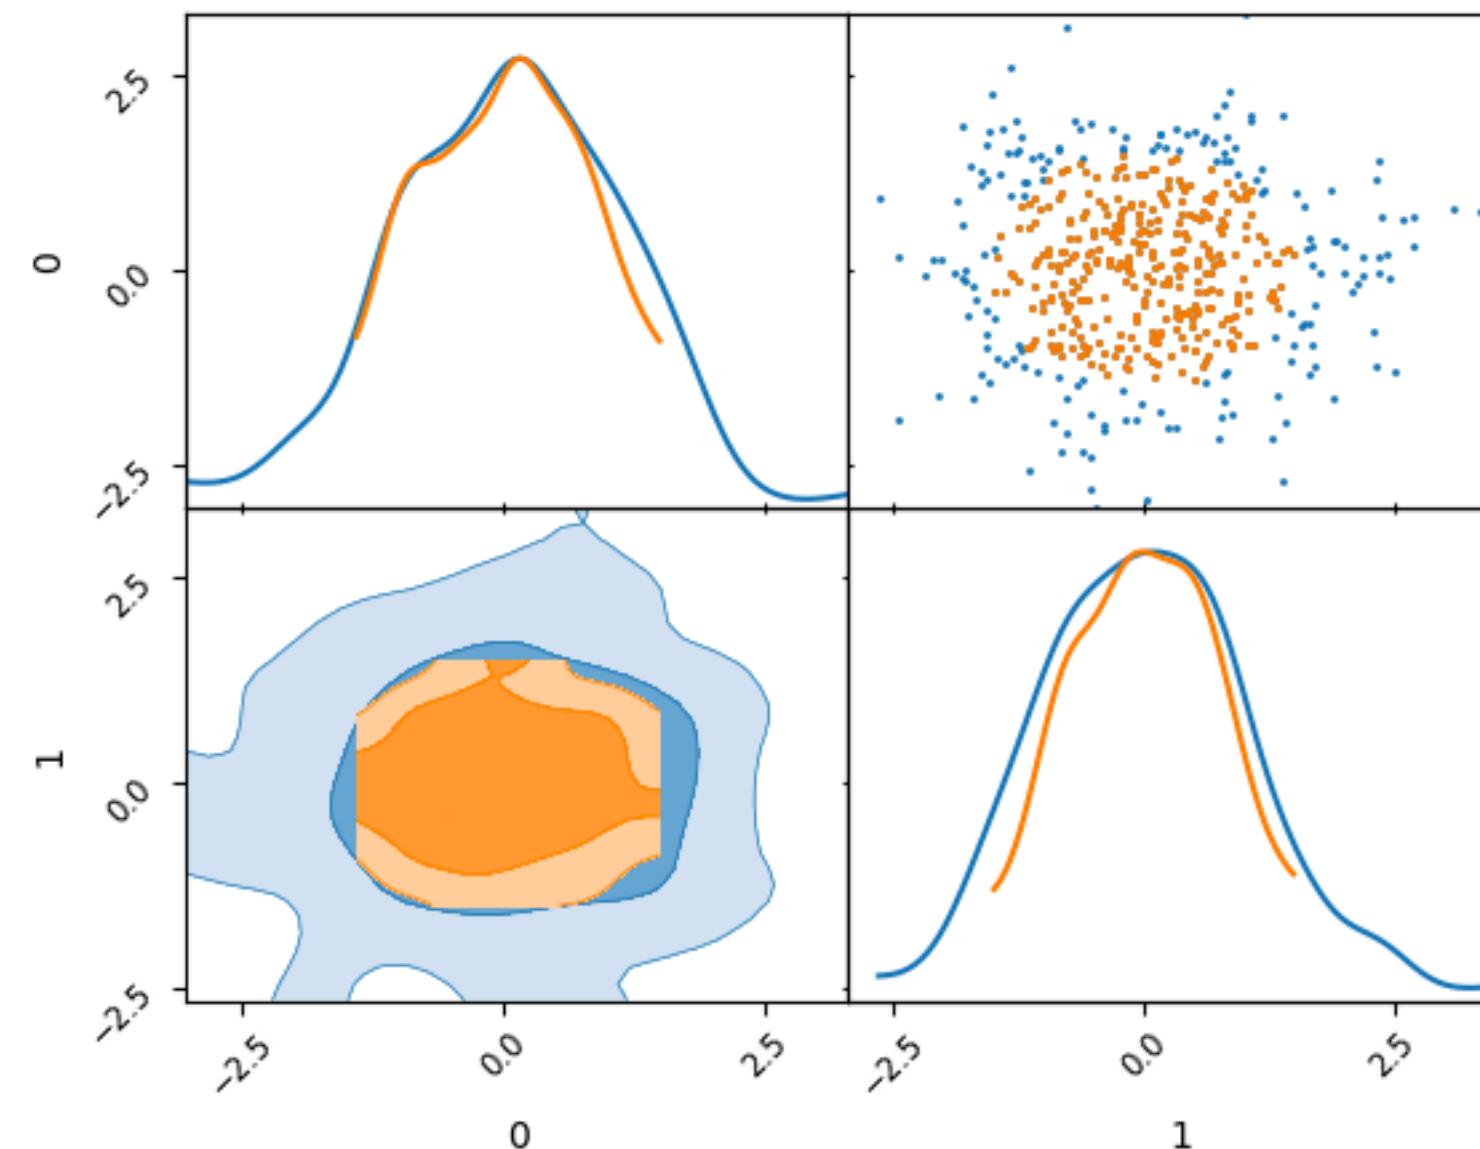


# Modelling distributions



# Nested Sampling with distributions

## “Rejection sampling”

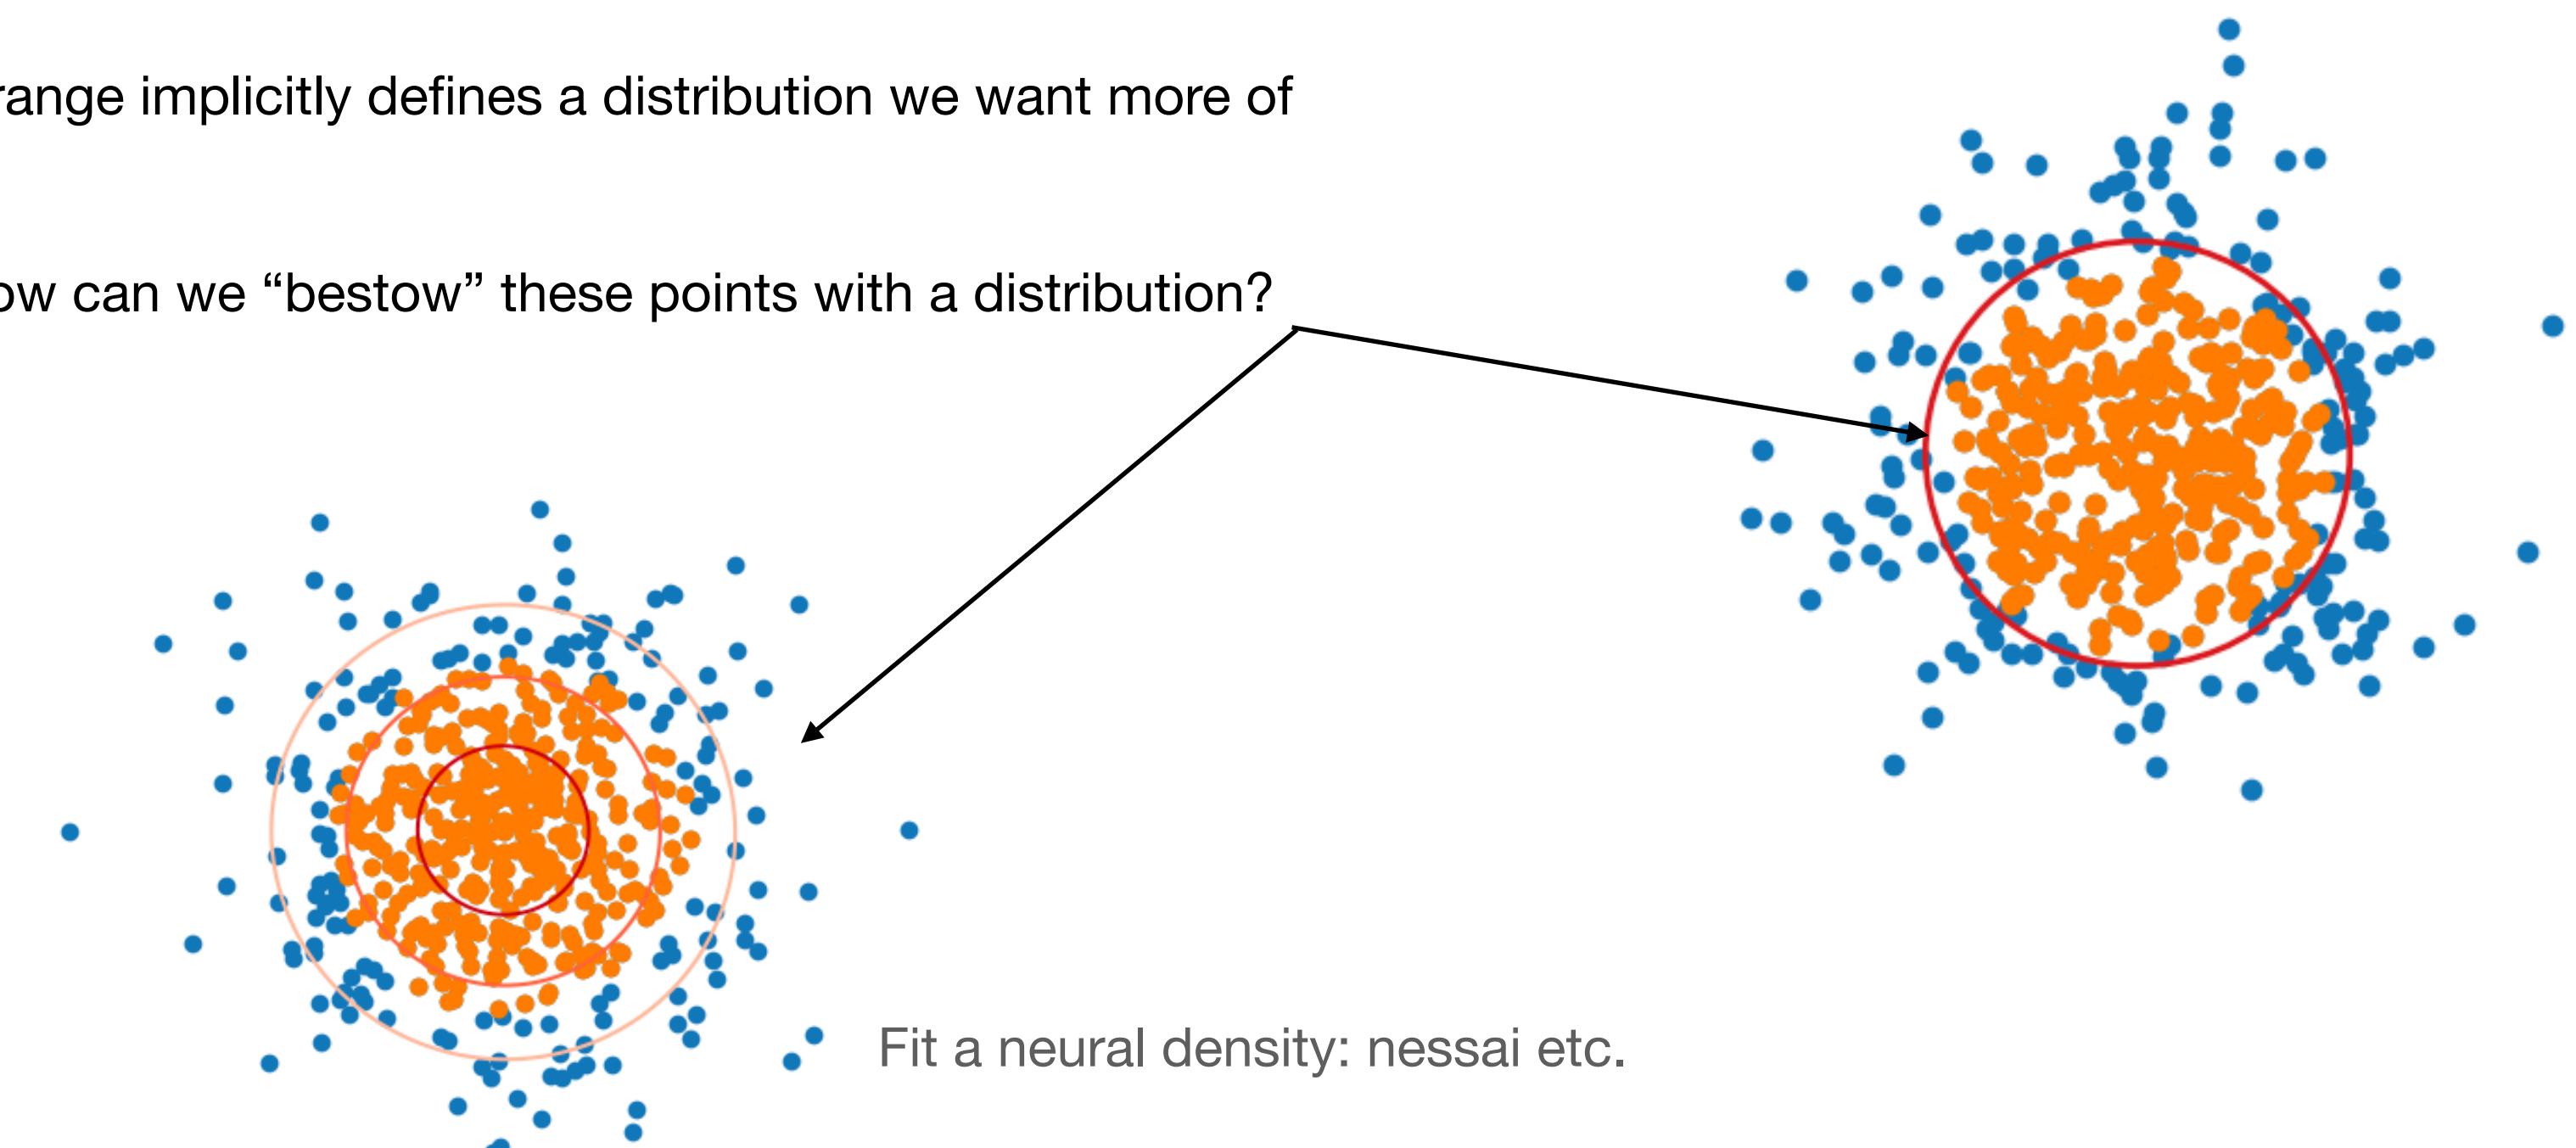


- Blue Samples from a known Distribution
- Orange Samples from likelihood constrained region

Orange implicitly defines a distribution we want more of

How can we “bestow” these points with a distribution?

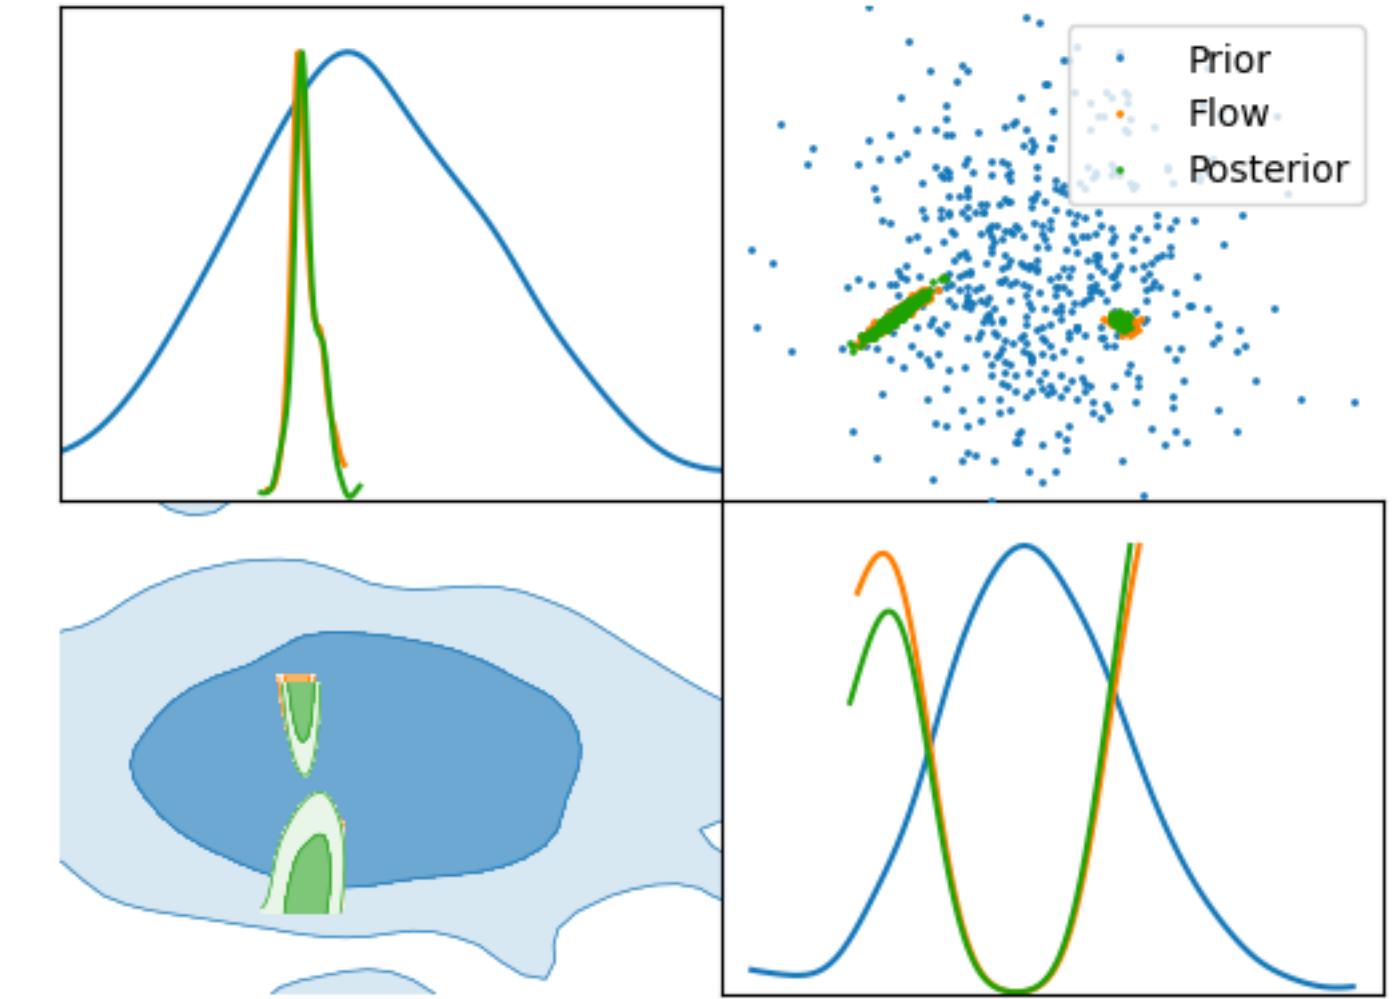
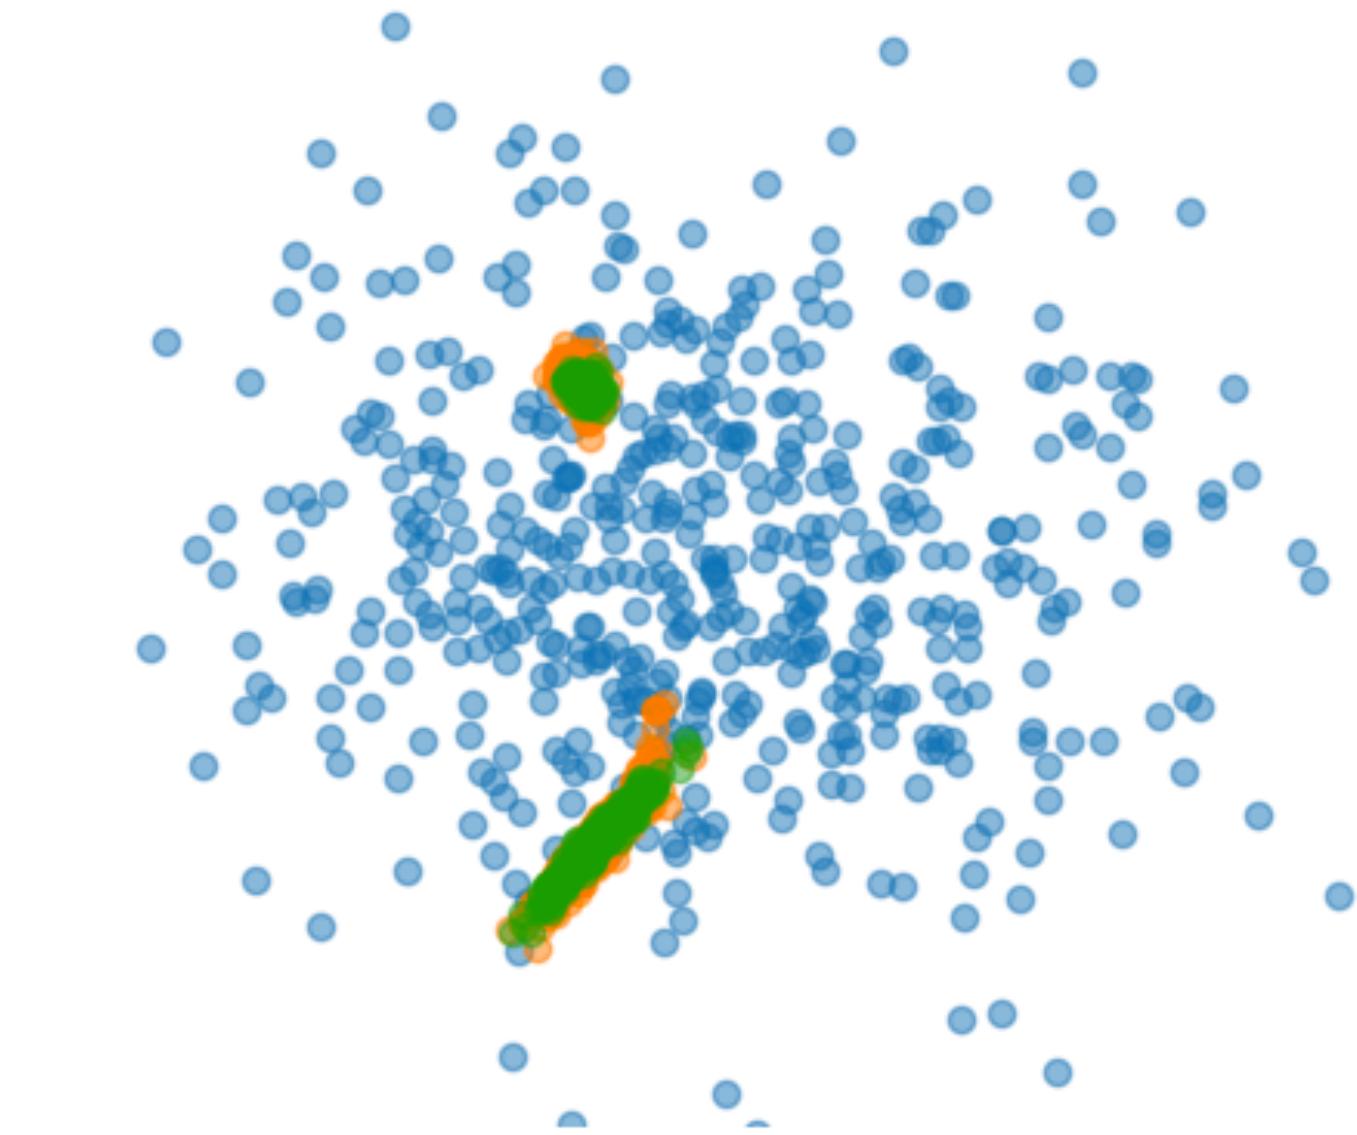
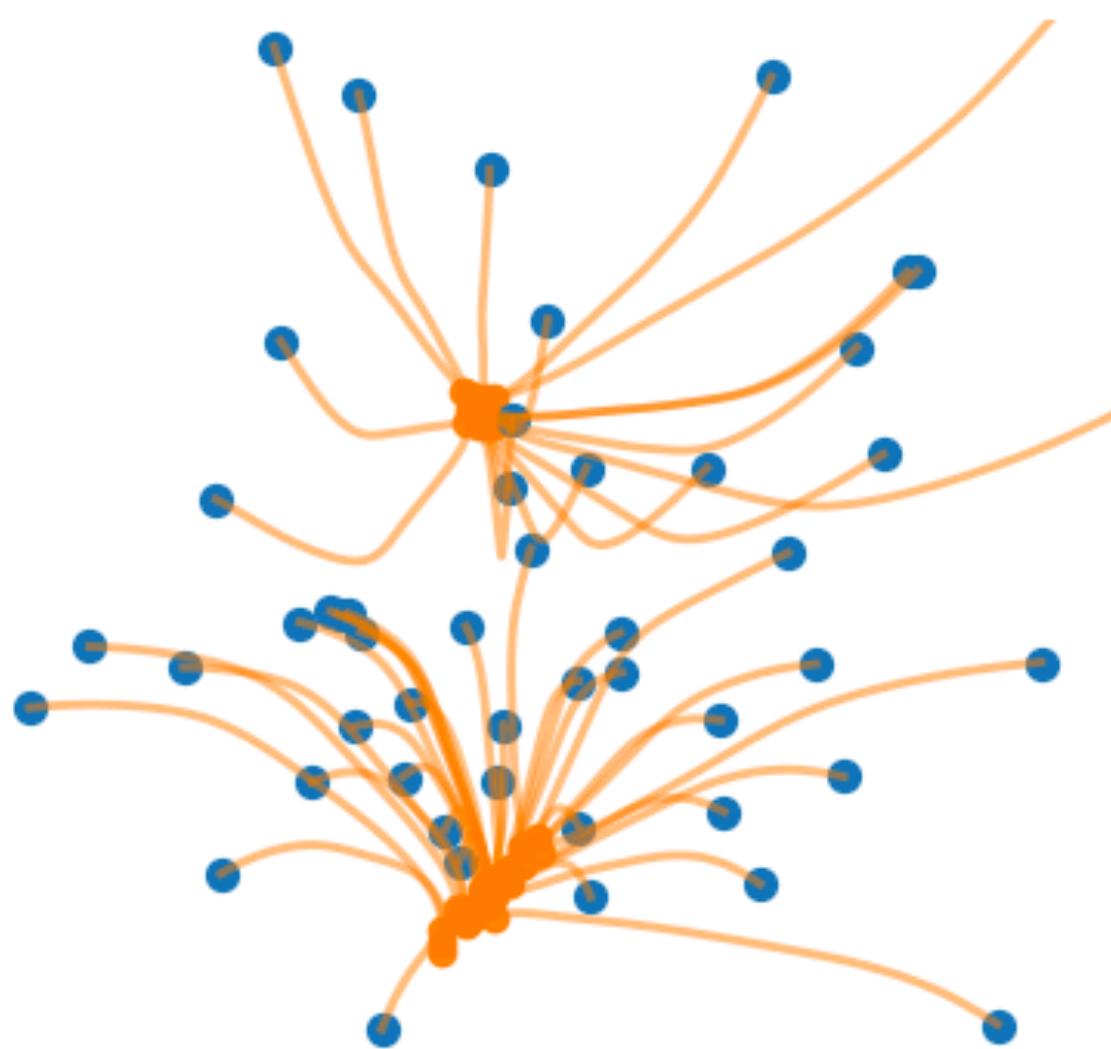
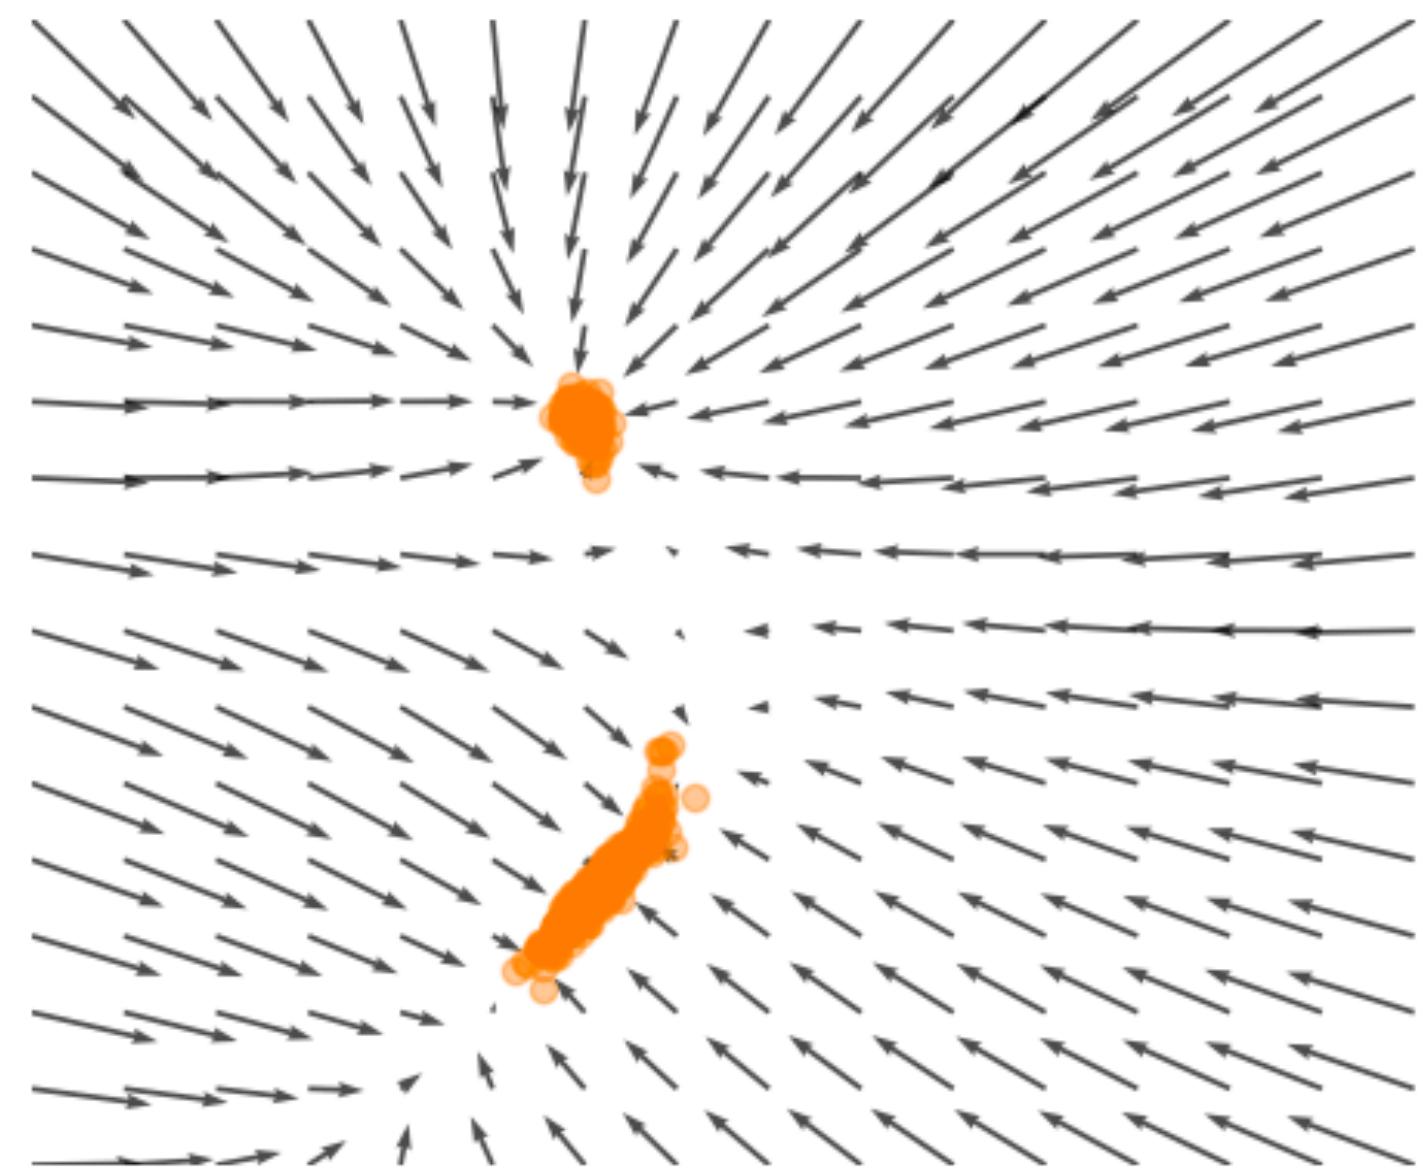
Fit an ellipse: MultiNest



# Diffusion Models

## The current vogue in density estimation

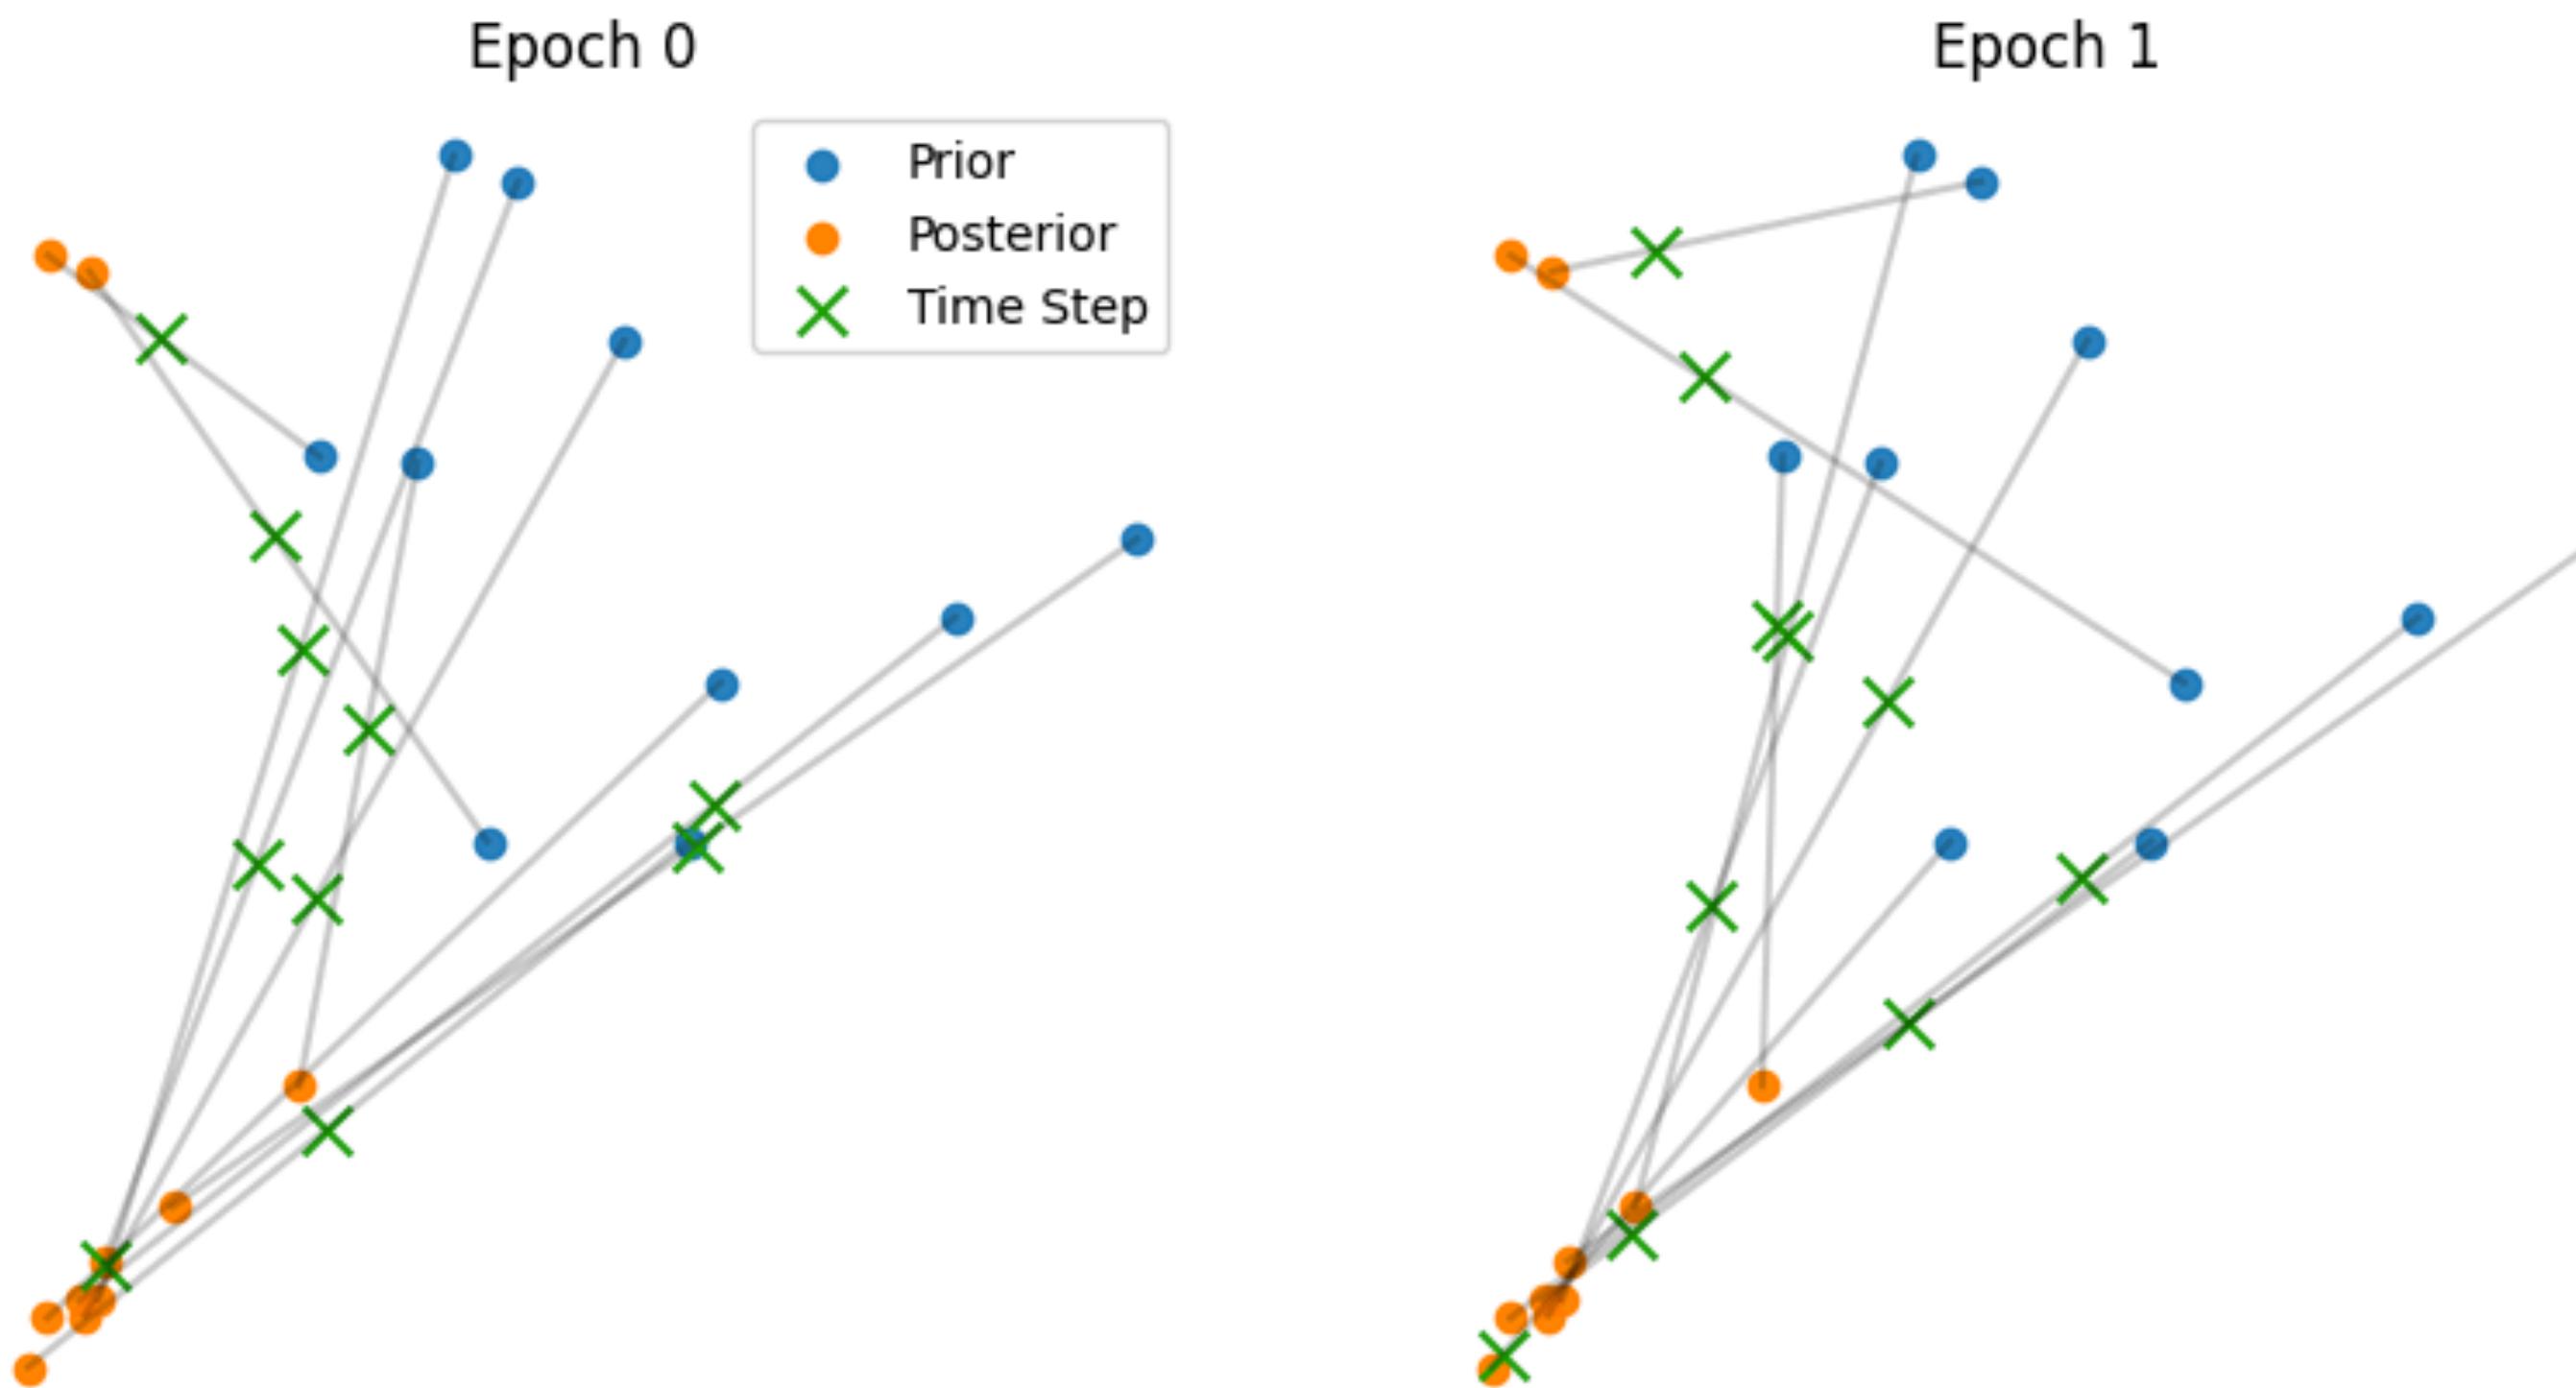
Optimal Transport, Diffusive processes, Continuous flows, Score based generative models, flow matching etc.etc. All the same thing!



Neural Networks Learn gradient fields (Score) that map one distribution to another

# How it works

[2302.00482] Flow matching, incredibly simple objective!



```
@partial(jit, static_argnums=[0])
def loss(self, params, batch, batch_prior, batch_stats, rng):
    """Loss function for training the CFM score.

    Args:
        params (jnp.ndarray): Parameters of the model.
        batch (jnp.ndarray): Target batch.
        batch_prior (jnp.ndarray): Prior batch.
        batch_stats (Any): Batch statistics (batchnorm running totals).
        rng: Jax Random number generator key.

    """
    # sigma_noise = 1e-3
    rng, step_rng = random.split(rng)
    N_batch = batch.shape[0]

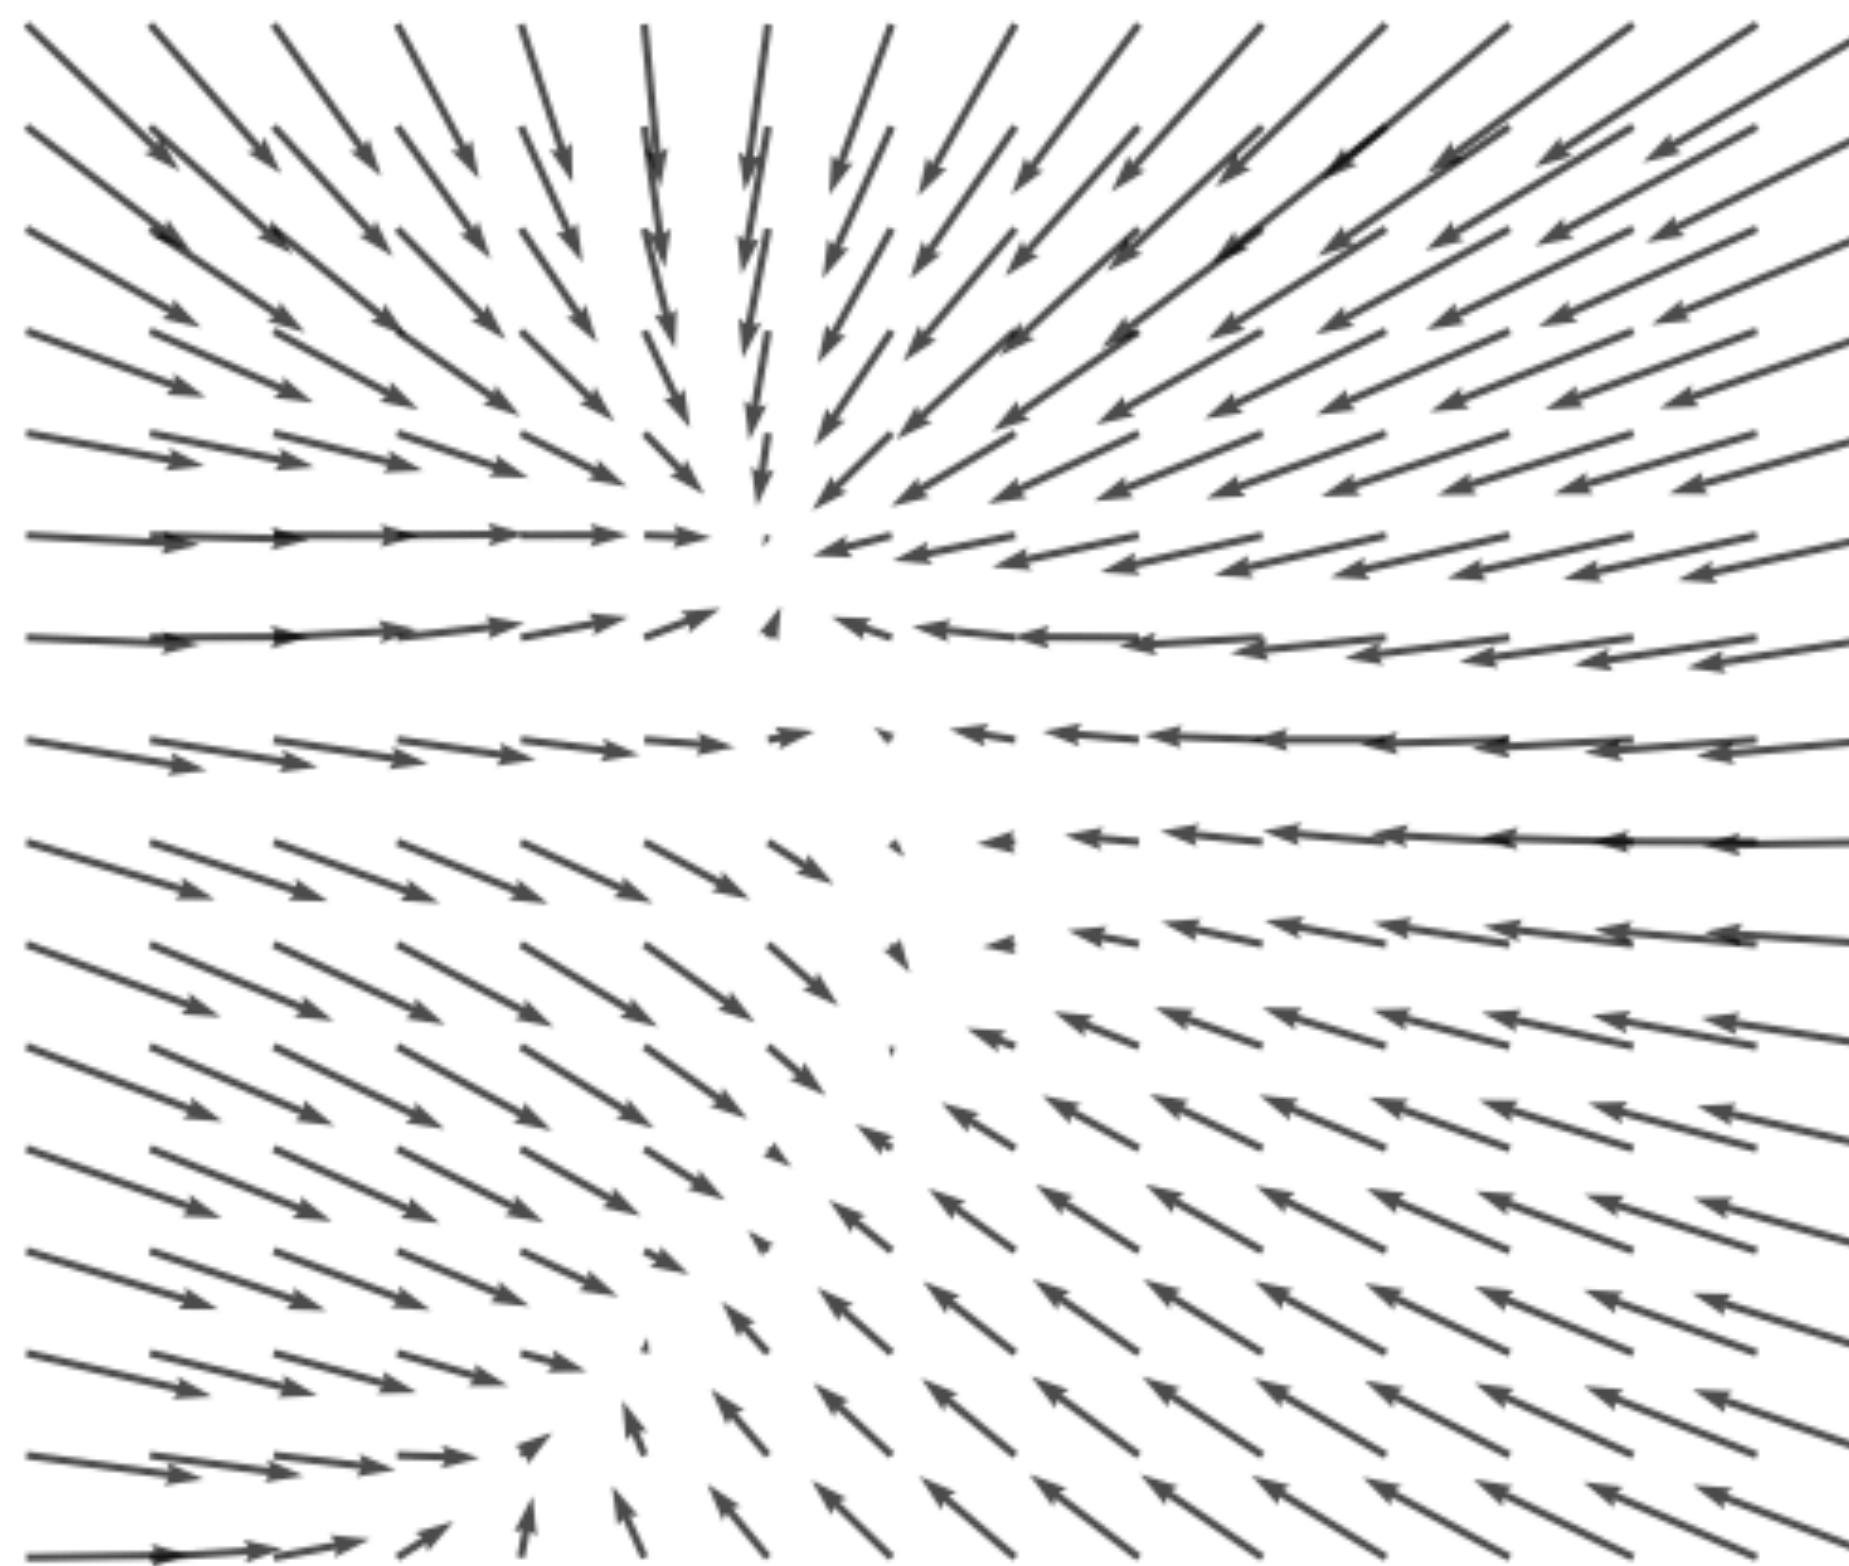
    t = random.uniform(step_rng, (N_batch, 1))
    noise = random.normal(step_rng, (N_batch, self.ndims))
    psi_0 = t * batch + (1 - t) * batch_prior + self.noise * noise

    output, updates = self.state.apply_fn(
        {"params": params, "batch_stats": batch_stats},
        psi_0,
        t,
        train=True,
        mutable=["batch_stats"],
    )
    psi = batch_prior - batch
    loss = jnp.mean((output - psi) ** 2)
    return loss, updates
```

- Pair up a random set of prior and target samples
- Generate a random time step for each pair,  $t \leftarrow [0,1]$
- Simple MSE regression on the prior-target vectors and a neural network learning the vector field at the time step

By shuffling and simulating a new time step we build up coverage of the whole space and trace out all paths

# What to do with a learned vector field



## Diffusion:

$$q(x_t | x_{t-1}) = \mathcal{N}(x_t \sqrt{1 - \beta_t} x_{t-1}, \beta_t I)$$

$$q(x_{1:T} | x_0) = \prod_{t=1}^T q(x_t | x_{t-1})$$

Solve a SDE, easy to simulate solutions from  
[\[yang-song.net/blog/2021/score\]](https://yang-song.net/blog/2021/score)

## Continuous Flows:

$$\frac{dy}{dt} = \nabla_\theta f(t, y(t)), \quad y(0) = y_0$$

Solve an ODE, slightly harder but nicer properties for science  
[\[2202.02435\]](https://arxiv.org/abs/2202.02435)

# Jacobians are an interesting part of this

## Often for scientific applications need this

